How Verified (or Tested) is My Code?
Falsification-Driven Verification and Testing

Alex Groce · Iftekhar Ahmed · Carlos Jensen · Paul E. McKenney · Josie Holmes

Abstract Formal verification has advanced to the point that developers can verify the correctness of small, critical modules. Unfortunately, despite considerable efforts, determining if a "verification" verifies what the author intends is still difficult. Previous approaches are difficult to understand and often limited in applicability. Developers need verification coverage in terms of the software they are verifying, not model checking diagnostics. We propose a methodology to allow developers to determine (and correct) what it is that they have verified, and tools to support that methodology. Our basic approach is based on a novel variation of mutation analysis and the idea of verification driven by falsification. We use the CBMC model checker to show that this approach is applicable not only to simple data structures and sorting routines, and verification of a routine in Mozilla's JavaScript engine, but to understanding an ongoing effort to verify the Linux kernel Read-Copy-Update (RCU) mechanism. Moreover, we show that despite the probabilistic nature of random testing and the tendency to incompleteness of testing as opposed to verification, the same technqiues, with suitable modifications, apply to testing as well as to formal verification. In essence, it is the number of surviving mutants that drives the scalability of our methods, not the underlying method for detecting faults in a program. From the point of view of a Popperian analysis where an unkilled mutant is a weakness in a "scientific theory" of program behavior, it is only the number of weaknesses to be examined by a user that is important.

1 Introduction

> "Every 'good' scientific theory is a prohibition: it forbids certain things to happen. The more a theory forbids, the better it is."
> -Popper, *Conjectures and Refutations: The Growth of Scientific Knowledge* [80]

Software model checking [25] has recently, thanks to improvements in model checking tools (and advanced SAT/SMT solvers), become a potentially valuable

Alex Groce
Northern Arizona University

```
#include "sort.h"
int a[SIZE];
int ref[SIZE];
int nondet_int();
int main () {
  int i, v, prev;
  int s = nondet_int();
  __CPROVER_assume((s > 0) && (s <=SIZE));
  for (i = 0; i < s; i++) {
    v = nondet_int();
    printf ("LOG: ref[%d] = %d\n", i, v);
    ref[i] = v; a[i] = v;
  }
  sort(a, s);
  prev = a[0];
  for (i = 0; i < s; i++) {
    printf ("LOG: a[%d] = %d\n", i, a[i]);
    assert (a[i] >= prev);
    prev = a[i];
  }
}
```

Fig. 1 CBMC harness to check a sorting routine.

tool for developers of critical software modules who want to either perform a very aggressive search for bugs or, ideally, prove correctness of their code. Tools such as CBMC [57] (the C Bounded Model Checker) allow a software engineer to model check code by writing what is essentially a generalized test harness [41,45][1] in the language of the Software Under Test (SUT). Figure 1 shows an example CBMC harness for sorting routines. Only a few aspects differ from normal testing. First, nondet_int in CBMC can return any value. It is not equivalent to a "random" choice but true nondeterminism: CBMC will explore all values of the type. The __CPROVER_assume statement has the usual assume semantics [29,44]: CBMC ignores all executions that violate assumptions.

CBMC compiles a harness and the SUT (here a quicksort implementation) into a goto-program, instruments this program with property checks for assertions, array bounds violations, etc., and then unrolls loops based on a user-provided unwinding bound to produce a SAT problem or SMT constraint such that satisfying assignments are representations of a trace demonstrating a property violation, known as a counterexample [24]. For CBMC, this means that if *any possible execution allowed by the harness* violates any properties checked, a counterexample will be produced. This includes user-specified assertions and automatically generated properties such as array bounds and pointer validity checks. One generated property is that no loop in the program executes more than the *unwinding bound* times. For example, if we run CBMC on the harness shown and set the unwinding bound to 3 and add -DSIZE=2, we will check the correctness of the SUT over all possible arrays of size 2 or less, including checking that sorting never requires passing through any loop more than 3 times.

When a model checker produces a counterexample, a developer's task is straightforward, if sometimes difficult: either the SUT has a fault, or the harness itself is

---

[1] By a harness we mean a program that defines an environment and the form of valid tests, and provides correctness properties.

flawed. In both cases, the output of the verification effort is the counterexample trace, which is full of evidence as to the reason for the failure to verify the SUT. Moreover, any solution (fix to SUT or harness) is easily checked: if it is correct, the model checker stops reporting the previous counterexample.

Unfortunately, model checkers do not invariably report counterexamples: eventually the SUT is likely to satisfy the properties encoded in the harness! It is in this case that problems arise: what, precisely, has been verified? Is the SUT actually correct? Formal verification is not only subject to the problems that make "no faults detected" results dubious in testing [38,40], but also to more subtle problems. For example, an incorrect `assume` statement may constrain a system so that not only are there no counterexamples, there are no (interesting) executions of the system at all. Moreover, formal verification tools are themselves extremely complex software artifacts, and, like production compilers [89], may themselves have serious bugs that produce wrong results [26]. In the course of this research, we have ourselves encountered several tool-induced incorrect verifications.

The problem of checking verification results has concerned the model checking community for some time, and resulted in efforts to define *coverage metrics* for model checking [53,23]. While such metrics are interesting and useful, they have typically been aimed at hardware verification, and most useful to experts in formal verification. In this paper, we adapt traditional mutation testing [18,63] to the problem of software verification. A mutant of a program is a version of the program with a small syntactic change. The idea behind mutation testing is that a good test suite will be able to detect when (as is usually the case) such a change introduces a bug in the SUT. In the case of bounded model checking, since we aim at (bounded) verification rather than merely good testing, surviving mutants are likely to indicate a real problem.

In software engineering research, mutation is often used only as a way to compare competing test suites by comparing kill rates [34,35]. This is not enough for verification. The typically small scope of the code to be verified, and the presumed importance of verified code suggests an approach in which *individual mutants* are examined by the developer. Without additional assistance, such an approach cannot scale. This paper aims to describe how to make this seemingly too-demanding approach practical for real verification tasks.

Our basic idea is to use mutants *throughout the verification effort*, even in choosing a bound for bounded model checking. At each stage the developer examines the currently surviving mutants, either by inspecting the mutated code or (when this does not make the reason the mutant is not detected clear) looking at *successful executions covering the mutant but satisfying the specification given in the harness*. For critical verification tasks, we suggest that developers not only examine the passing executions of surviving mutants, but the passing executions of *killed mutants*. While examining test cases that do not kill a given mutant could be useful in traditional testing, the model checker makes a much more potent investigation possible, where a developer can constrain the behavior to force the mutant's behavior to matter, if that is possible, and automatically find passing executions that maximize coverage (including the mutated code). We also propose that a developer should use mutants of the test harness itself to ensure that no similar harness has a better mutant kill rate, and that most mutants of the harness reject the SUT itself.

1.1 Contributions

This paper is an extension of a paper presented at the 30th IEEE/ACM International Conference on Automated Software Engineering in 2015 [39]. The core contribution of that earlier paper was a *falsification-driven* verification methodology using mutants to aid developers in understanding "successful" verification results, determining when a harness is flawed, and correcting the harness. It showed how to use mutation testing to choose a problem size in bounded model checking, how to mutate a harness to determine if any similar harnesses have an equal (or better) mutation kill rate, and most importantly, how to modify CBMC, a harness, and mutants to automatically produce successful high-coverage executions covering mutated code in order to understand mutant (and thus harness) behavior. This approach, unlike a simpler method of searching for cases where the mutated and original code behave differently for identical inputs, in principle applied to verification of reactive and concurrent systems, where there is no simple notion of identical inputs. It also proposed the use of mutation analysis to gain limited confidence in program correctness even past model checker scalability limits. At a more general level, the original paper discussed the fundamental nature of "verification" in a real-world context where specifications are never known to be complete. The central concept of that paper, and this extension, is that falsification, as in Karl Popper's famous philosophy of scientific discovery [79], is the critical element of efforts to understand systems, efforts that are always provisional by nature: e.g., most program verification (and even more so, testing) efforts. Popper's approach therefore forms a useful conceptual framework for verification efforts: rather than focusing on what can be proven about a program, we propose that correctness-determination efforts focus on how a verification distinguishes the "real" program from similar alternative programs that do *not* match the theory of program behavior. Such an approach still aims at verification as a final outcome, but continually evaluates and refines that verification effort by its ability to *falsify* rather than to verify.

In this paper, we extend the contributions of the previous work by further elaborating the approach to bounded model checking presented there, but, most significantly, we extend the same ideas to automated software testing, where there are new challenges and it is less clear that the underlying concepts are sound. In our previous work, we dismissed the possibility of using our approach for testing, because when a typical test generation approach fails to find a failing test (that is, testing's version of a counterexample), it does not mean the property is insufficient, or even that the generation is weak. Testing is usually essentially probabilistic. Here we show that this limitation is not fundamental, and the same principles can apply to serious automated test generation efforts as well. The general ideas are analogous, including the modification of a test generation tool to generate high-coverage tests that 1) fail to kill a mutant but 2) cover the mutated code, in order to facilitate understanding of the limits of a harness, but the details require considerable modification in order to adapt to the probabilistic context of test generation. For example, the equivalent of a problem size for bounded model checking is a test budget for automated test generation; however, rather than a fixed outcome, the problem in this case is to present a likelihood curve to users, and allow them to make tradeoffs based on that curve.

```
   #include "sort.h"
   void quickSort( int a[], int l, int r)
   {
     printf ("LOG: called with l=%d, r=%d\n", l, r);
     int j;
 9 if( l < r )
      {
        // divide and conquer
        j = partition( a, l, r);
        quickSort( a, l, j-1);
        quickSort( a, j+1, r);
      }
   }
   int partition( int a[], int l, int r) {
     int pivot, i, j, t;
     pivot = a[l];
     i = l; j = r+1;
26 while( 1 )
      {
28      do ++i; while( i <= r && a[i] <= pivot );
        do --j; while( a[j] > pivot );
30      if( i >= j ) break;
31      t = a[i]; a[i] = a[j]; a[j] = t;
      }
     t = a[l]; a[l] = a[j]; a[j] = t;
     return j;
   }
   void sort(int a[], unsigned int size) {
     quickSort(a, 0, size-1);
   }
```

Fig. 2 Quicksort code.

We still initially present our approach in the context of formal verification, where it is most obviously useful, and where there is a higher chance of killing enough mutants to make manual analysis of un-killed mutants feasible.

2 A Simple Example Verification

As an example of the proposed verification methodology, consider again the harness shown in Figure 1. If we take an early Google result for "quicksort in C" [61], shown in Figure 2[2], we can model check it using the harness, defining SIZE=2 and setting the unwinding bound to 3 (we need one more unwinding than the maximum number of items in the array). CBMC reports VERIFICATION SUCCESSFUL in less than a second. Have we verified what we want to verify?

2.1 Finding a Good Problem Size

The first question we face is whether 2 is a good maximum array size to examine. The problem of determining a completeness threshold (an execution-length bound

---

[2] In fact, that actual code is incorrect, with an access a[i] that does not properly use short circuiting logical operators to protect array bounds; CBMC detected this, and we fixed it for this paper.

sufficient to prove correctness in all cases for a given property) for bounded model checking is fundamentally difficult [58] and is, for real-world C programs, more an art than a science at present[3]. Are there bugs for which 2 is too small an array size? In order to find out, we generate a set of mutants for `quicksort.c`. Using the mutation tool for C code developed by Jamie Andrews [5], we can produce 81 mutants of this code in less than a second. We then run the harness with unwinding bound 2 (and `SIZE=1`) on each of the 81 mutants. The process takes less than a minute and a half (on a MacBook Pro with 16GB RAM and dual-core 3.1GHz Intel Core i7, using only one core). CBMC reports that 6 mutants do not compile (these remove variable declarations, for the most part), 4 are detected by the harness (a counterexample is produced: we say the mutant is killed), and 71 mutants pass without detection (the verification is successful, in which case we say the mutant survives). Clearly length 1 arrays are not sufficient to detect even the most glaring bugs in a sort algorithm (no surprise: all size 1 arrays are sorted). What about our choice of size 2? Re-checking the mutants with this bound (dropping those already killed by the smaller bound) takes slightly over 13 minutes (due to one mutant requiring over 8 minutes to model check) and reduces the number of surviving mutants to 26. We could inspect these mutants by hand, but it seems highly unlikely that a complete verification over all possible arrays with a good specification for sorting would produce such a poor mutation kill rate. If we increase the size limit to 3 and unwinding 4 (now the analysis takes just over 33 minutes), only 8 mutants survive. Note that each problem, due to the harness' assignment of `s` to any size smaller than the current size, includes all smaller problem sizes. This makes the behavior of the verification problem size setting match that of CBMC where an unwinding bound is a maximum, rather than a fixed size. We assume inclusiveness in this paper[4].

At this point, we can increase `SIZE` to 4 (which will kill one additional mutant), but the time required to check the remaining mutants is growing rapidly. In fact, completing the check for size 4, even though only the original program and 8 mutants have to be checked, requires nearly 9 hours. When the model checking difficulty grows more slowly with problem size, we propose the simple (if highly imprecise) heuristic of increasing size until the number of mutants killed does not increase with a step up in size (we call such a size *mutant-stable*). However, in many cases, such as this one, the time required to check mutants starts growing unacceptably. We propose a more efficient algorithm for finding a mutant-stable size below (Figure 7), and mutations can be checked in parallel, but the fundamental problem for size 4 (and above) is that some individual mutants require hours to model check. What is a developer to do?

2.2 Examining Surviving Mutants

The developer should turn to the surviving mutants. The 8 surviving mutants for size 3 are shown in Figure 3. The comment indicates the type of mutant, and the line number in the quicksort file is also given for reference. The relevant lines are marked in Figure 2. Some of these mutants are easily seen to be equivalent

---

[3] In our own practice, the most common way of setting it is to guess a bound and see if the resulting problem is too large for the available resources.

[4] There is one noted exception in Section 4.4.

```
 9 :   /*(rep_op)*/ if (l <= r)
26 :   /*(rep_const)*/ while(-1)
26 :   /*(rep_const)*/ while( ((1)+1))
28 :   /*(rep_op)*/ do ++i; while(i<r && a[i]<=pivot);
28 :   /*(rep_op)*/ do ++i; while(i!=r && a[i]<=pivot);
28 :   /*(rep_op)*/ do ++i; while(i<=r && a[i]<pivot);
30 :   /*(rep_op)*/ if( i > j ) break;
31 :   /*(del_stmt)*/ t=a[i]; /*  a[i]=a[j]; */  a[j]=t;
```

Fig. 3 Surviving mutants at `SIZE=3`.

```
LOG: ref[0] = 2147414872
LOG: ref[1] = 2147480408
LOG: ref[2] = -1073743560
LOG: called with l=0, r=2
LOG: called with l=0, r=-1
LOG: called with l=1, r=2
LOG: called with l=1, r=1
LOG: called with l=3, r=2
LOG: a[0] = 2147414872
LOG: a[1] = 2147480408
LOG: a[2] = 2147480408
```

Fig. 4 Witness to the harness' inability to kill the `del_stmt` mutant.

to the original code. For example, the two `rep_const` mutations simply change a `while(1)` into an equivalent infinite loop with a different constant non-zero value. These two mutants could in fact have been automatically removed from the set, like uncompilable mutants, by checking their compiled code for equivalence with the original program [78]. We suggest always pruning mutants via Trivial Compiler Equivalence (TCE). The remaining 6 mutants produce different binaries when compiled with an optimizing compiler, so require manual analysis. The 5 `rep_op` mutations all alter comparison operators by changing their value on one corner case, and we may suspect that quicksort is robust to, for example, changing `i <= r` to `i != r` since `i` is initially set to `l`, which we know to be less than `r`.

The `del_stmt` mutant, however, is clearly problematic. How can quicksort be correct if the inner loop's swapping of `a[i]` and `a[j]` is changed to instead copy `a[i]` to `a[j]`? The consequences of this mutant are clearly drastic, but why are they not detected by our harness? We find out by asking CBMC to produce an execution such that 1) the mutated code is covered 2) other coverage is maximized (to avoid degenerate executions, e.g., over size 1 arrays) and 3) the execution is not a counterexample. We have modified CBMC, and written instrumentation tools that produce a modified mutant and harness, allowing us to pose such queries (see Section 3). Running CBMC in this mode, with the target of maximum branch coverage and statement coverage of the `del_stmt` mutant (actually the statement after it, since it no longer exists), we produce the witness in Figure 4 in less than a minute[5]. Our harness checks that the array `a` is sorted after the call to `sort`, but it does not check that it is a permutation of the input!

We might have discovered this problem by a different method: if we remove the call to `sort` in the harness, and replace it by a loop assigning `nondet_int` to every

---

[5] We show the output of the print statements, not the full CBMC trace: this is what a developer will examine first.

```
...
   int i, v, count, qcount, prev;
...
  sort(a, s);
  // Pick a value to check
  v = nondet_int();
  count = 0;
  qcount = 0;
...
  for (i = 0; i < s; i++) {
...
     if (ref[i] == v)
       count++;
     if (a[i] == v)
       qcount++;
...
   assert (count == qcount);
}
```

Fig. 5 Modifying the harness to ensure `a` is a permutation of `ref`.

element in array `a` (a kind of most-general any-order type-correct "mutant" of `sort`), we can run the modified CBMC and see examples of executions our harness allows, which should include any sorted array. The problem with this method is that, while it sometimes works, CBMC is also free to set all elements in all arrays to 0, and to generally provide an uninformative example of a successful execution. The requirement to cover a mutant (and as much other code as possible) helps guide CBMC to a successful execution that is likely to be incorrect, because a non-equivalent mutant changes the original program's behavior. Moreover, while the problem with the harness in this case is simple, understanding arbitrary "passing" but wrong executions can be very difficult without the ability to think about a specific bug the model checker is missing. Moreover, basing the production of witnesses on mutants allows us to compare harnesses even over killed mutants: if one harness reduces the set of passing executions for a mutant, it is arguably a better verification of correctness than one allowing more executions of the mutant, even if both produce a counterexample killing the mutant. Unlike traditional mutation analysis, we can take the question "how killed is this mutant?" seriously because we aim at exhaustive testing. A harness is most effective with respect to a mutant if it allows no executions covering the mutant to pass.

The witness tells us that the sorting harness is too weak. We say that a harness is weak if it fails to detect incorrect executions. One harness is stronger than another if it detects more failures; we can indirectly estimate strength by determining how many mutants a harness can kill at a given problem size, and how executions covering killed mutants can still satisfy the harness. Figure 5 shows how to modify the sorting harness to check for permutations[6]. Because CBMC is exhaustive, instead of performing a complete check for permutation, we can detect violation of the property by letting `v` be any value, and ensuring that both `a` and `ref` contain the same number of elements equal to `v`. If `a` is not a permutation of `ref`, there exists a `v` such that this is not true, and we can rely on CBMC to report

---

[6] In fact, if we choose a `val` to check before we assign to `ref`, we could completely dispense with storing `ref` at all.

it as part of a counterexample. While a CBMC harness resembles a program to test the SUT, it can make use of unusual specifications relying on exhaustiveness.

If we modify the harness as shown, we can re-check our mutants (including those TCE would remove). With the revised harness, checking mutants at `SIZE=1` takes slightly longer (8 more seconds) and kills the same number of mutants, since the problem is the size, not the harness. At `SIZE=2` mutant kill results are again unchanged, but analysis now completes in about 5 minutes. Finally, at `SIZE=3`, we kill the `del_stmt` mutant that previously survived, after only 14 minutes, not much longer than at `SIZE` 2 with the weaker harness. The `SIZE` 3 verification is stable. Checking stability by running `SIZE` 4 now only requires slightly more than an hour, nearly an order of magnitude faster than before.

As briefly mentioned in the introduction, it is also possible to understand a mutant by modifying the harness to call both the mutated code and the original code on the same inputs, and search for witnesses where 1) the execution is passing but 2) the return value(s) for the mutant differ(s) from the original. However, this increases the complexity of the model checking problem (checking equivalence of two functions is often harder than specifying valid executions) and does not easily apply to any verifications other than simple function calls. For example, forcing the same interleavings in threaded programs, or detecting all differences in state-modification for reactive code is often infeasible or requires significant human intervention. While we do apply differential checks in some cases below, we do not propose this as a core technique suitable for general-purpose falsification-driven verification.

2.3 Mutating the Harness

Previous efforts to understand model checking results have also considered mutants to the property, usually given as a temporal logic formula [13]. Once a developer is satisfied with a harness, has a mutant-stable bound for verification, and is convinced all surviving mutants are semantically equivalent to the original program (or, if not equivalent, also satisfy the same correct specification), we propose the developer mutate the test harness itself. The idea is to check that 1) most mutants of the harness reject the SUT and 2) the remaining mutants have a mutant kill rate no greater than that of the harness. For the fixed sort harness, there are 61 mutants, of which 2 do not compile. Of these, 40 produce an incorrect counterexample for the original, correct, quicksort. An additional 10 have mutant kill rates worse than the original harness (from as low as 5% of mutants killed to only a few percent worse than the fixed harness). The remaining 9 harness mutants have the same ability to kill mutants as the original harness. Most of these involve modifying a relational operator in a loop or an assumption in a way that preserves semantics. The only interesting surviving harness mutant is one that removes the assignment of a fresh non-deterministic value to `v` after the call to `sort`. This means the check for permutation difference will always be performed on the last element of `ref`. On reflection, it seems plausible that this is sufficient to produce a counterexample for all the quicksort mutants, but it is clearly not an improvement to the harness, in terms of either verification strength or clarity.

In addition to showing the current harness is at least a "local minima" with respect to mutants, mutation analysis of the harness also provides some evidence

of our technique's ability to detect subtle specification and environment flaws. In particular, it shows the value of inspecting all surviving mutants. One mutant modifies the assumption on `s` to be `s < SIZE` rather than `s <= SIZE`, which is the same as lowering the `SIZE` by one; this is a fairly easy mistake to make in a harness (or any code). This reduces the effectiveness of the verification by 19 mutants, so is likely not to escape notice, and would also (in our framework) simply result in a higher size being chosen as mutant-stable. Deleting the assignment `prev = a[i]`, however, only kills 4 fewer mutants than the original harness. Traditional coverage and some model checking coverages cannot detect this problem: because of the assignment to `prev` outside the loop, the variable *is* used in the specification, and in fact used to detect many faults (it eliminates any mutants that can cause `a[0]` to not be the least element). The harness "covers" all behavior of quicksort in general, since the permutation requirement remains in place. However, it cannot detect versions of quicksort that 1) preserve permutation and 2) make the first element correct, but 3) don't always sort the array correctly. In particular, the call to `quickSort` with `j + 1` can be removed or modified to `j + 2`. Examining the deleted/removed recursive calls shows the developer the problem in this case. Our modified CBMC easily produces a witness showing a permuted array with correct `a[0]` but with out-of-order later elements.

The basic flow of our approach, for bounded model checking, can be summarized as follows:

1. Generate mutant set $M = m_1 \ldots m_n$ for the program P.
2. Prune M into M' by equivalence classes based on optimizing compiler output, removing mutants that fail to compile or are equal to the original code.
3. Set unwinding depth/problem size U to 0.
4. Set $r = 0, r' = 1$.
5. While $r \neq r'$:
    (a) Set $U = U + 1$.
    (b) Set $r = r'$
    (c) Set $K = \emptyset, S = \emptyset$.
    (d) Check each mutant $m_i \in M$ using H and size U: if m_i is killed, $K = K \cup \{m_i\}$, otherwise $S = S \cup \{m_i\}$.
    (e) Set $r' = |K|/|M|$.
6. Examine each mutant in S. Remove those that are, by inspection, semantically equivalent to P.
7. Modify harness H for mutants in S that indicate a clear violation of the specification, easily understood, until H kills all such mutants. Remove them from S and add them to K.
8. For remaining mutants in S, generate a successful execution that covers the mutant but satisfied H. If the execution is degenerate, add constraints removing that class of execution until a witness to an incorrect, mutant-covering behavior is produced. Use this to modify H and move newly killed mutants from S to K.
9. Take mutants in $m_i \in K$, and check whether there exists a successful execution of m_i satisfying H. Examine and constrain each such execution to remove degenerate solutions, modifying H as needed.

10. Compute mutants M_H of the harness, and check that all mutants either: produce a counterexample for the original program P or have a kill rate $\leq$ the kill rate for H.

3 Algorithms and Techniques

Falsification-driven verification is a semi-automated approach that relies heavily on algorithmic and tool support. While the typically smaller scope of code targeted for verification (vs. testing) makes the work easier, it is not likely to be feasible without automation of many subtasks. Existing tools make producing a set of mutants and checking them using a harness relatively easy, but other tasks require new algorithms and tools. Figure 6 shows the basic flow, which is directed not by a fixed algorithm but by the intelligence (guided by experience) [84] of the developer. Novel tools or techniques are on the right side of the diagram (mutation analysis itself is not novel, but our tool for integrating this with the model checker and recording results for use by other parts of the tool-chain is non-standard), other than the model checker itself.

Figures 7-9 show core algorithms (implemented as prototype tools in our framework). In these algorithms $O(S)$ is a function mapping an abstract size into particular options, e.g. `-DSIZE`. The uses of these algorithms are described at a high level in the introductory example, and in the case studies below. One additional requirement is a version of CBMC capable of converting built-in assertions checks (e.g, bounds checks, pointer dereference, division by zero) to assumptions. For harness assumptions, this is done by automatic source-to-source transformation (Figure 8, procedure `cover-harness`), but CBMC's internal constraints have to be handled inside the model checker. We implemented a new CBMC command-line option, `--find-success` that provides this functionality. In all algorithms, `check` means running CBMC as usual, with any needed automatic property checks, while `scheck` indicates running CBMC with `find-success` enabled. In Figure 6 we assume the use of a modified version of CBMC.

The `find-size` algorithm (Figure 7) finds a suitable problem size and returns the set of surviving mutants for a harness and program, performing as few model checker calls as possible (once we know a bound is not stable, we move on to the next bound). This algorithm can be easily parallelized by running mutants in the `for` loop at the same time, with any `goto TOP` killing all CBMC runs not terminated. The `maxcover` algorithm (Figure 8) returns for a given mutant and harness, a witness program trace that 1) covers the mutant and 2) covers as much other code as possible (in terms of branch coverage), using the `cover-harness` and `cover-mutant` procedures to instrument harness and mutant; it proceeds by starting with a minimal constraint on coverage (the trace must cover the mutated code) and increases this bound by incrementing it to one more than the actual coverage of the last witness found, until the model checker can prove the coverage is impossible. Other strategies for maximal coverage are possible (trying maximal coverage first, and decreasing the required coverage as attempts fail) but this approach minimizes the number of model checker runs that will fail to produce a witness, which is critical for performance reasons (see Section 4.4). The `check-harness` algorithm (Figure 9) analyzes harness mutants, producing a report of 1) harness mutants that are killed (either they do not verify the SUT or they have worse kill rates
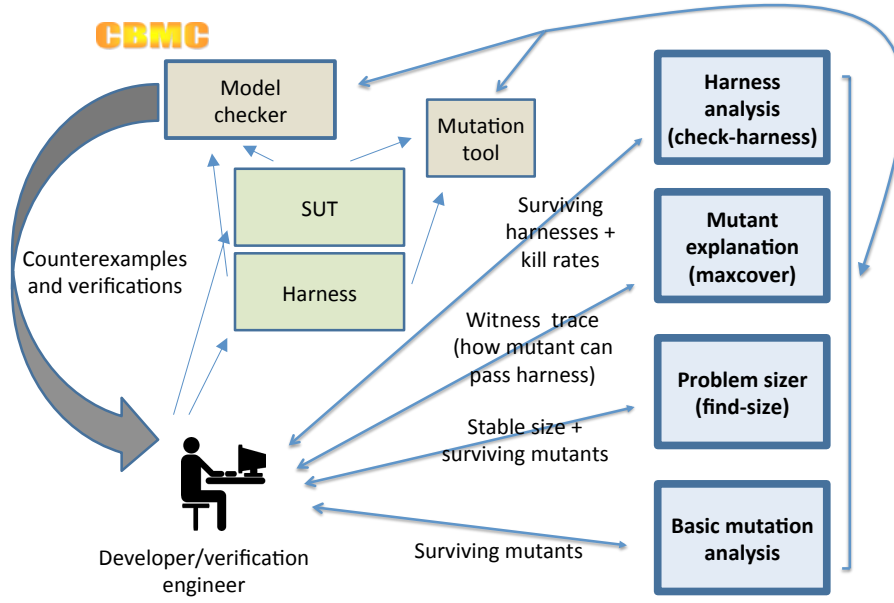
Fig. 6 Basic flow of falsification-driven verification.

than the original harness), that are equal to the original harness in strength, and that are stronger than the original harness. It also returns information on all mutants killed by any harness mutant (except those that reject the SUT) that are not killed by the original harness. The algorithm `killed`, not shown, simply returns the set of mutants killed by a given harness. In our implementations, these tools perform additional record-keeping. For example, harness analysis records killing counterexamples and execution times for each run. We also make use of convenience scripts such as a tool to automatically call `maxcover` on all mutants, which provides a measure of harness strength that is more fine-grained than a simple kill rate: harnesses can be compared by the maximum coverage of *all* mutants, even if they have the same kill rate. If one harness produces executions with lower coverage (or no executions at all) for some killed mutants, it is stronger. For some mutants, *any* passing executions show a harness flaw. While not polished enough for release, these tools (implemented as Python scripts) have proven robust in our experiments and are available, along with our experimental data and CBMC patch, at `https://github.com/agroce/cbmcmutate`.

```
(int, survivors) find-size (H, M, S₀: int,
                                   O: int → options,
                                   U: int → int)

S = S₀-1
r' = {} : mutant → result
TOP:
S = S + 1
r = r'
r' = {}
for m ∈ M:
    if m ∉ domain(r):
        r[m] = check(H,m,U(S),O(S))
        if r[m] == VERIFICATION FAILED:
          //once killed, assume always killed
          M = M \ m
    if r[m] == VERIFICATION SUCCESSFUL:
        r'[m] = check(H,m,U(S + 1),O(S + 1))
        if r'[m] == VERIFICATION FAILED:
          M = M \ m
          goto TOP
// No result changed, so S is mutant-stable
return (S-1, M)
```

Fig. 7 Algorithm 1: Finding size/unwinding bound and surviving mutants.

3.1 Adapting Falsification-Based Approaches to Automated Test Generation

3.1.1 Mutants and Manual vs. Automatically Generated Tests

In this paper, we do not propose the idea of using mutants to improve manually constructed test cases; that concept is essentially as old as the idea of mutation testing itself. Instead we focus on the testing equivalent of a verification harness: an automated test generation harness and the tests it produces. We believe that, despite superficial similarity, the two ideas are fundamentally different. In the first case, there is static, concrete, existing test suite. Mutation analysis is performed on this suite, and a developer writes concrete, specific, tests to kill mutants that are not equivalent. This may be a useful exercise, and may indeed somewhat improve the future defect detection powers of the suite [55, 1] to some extent. We suspect, however, that this method seldom immediately results in detection of an existing fault in the SUT. The developer of the test to kill the mutant is not, really, looking for a new way to test the system, but looking for a specific input that causes the mutant to behave incorrectly. Again, this may improve the suite, but seems unlikely to detect other, lurking, incorrect behaviors. To our knowledge, there are no human studies showing that such changes result in immediate fault detection.

However, the story is quite different with an automated test generation system based on a harness defining tests to be generated, for instance a sophisticated random test generator for a specific domain [42, 89, 51, 68]. In that setting, modifying the harness so that a test is generated that kills the mutant does more than kill a specific mutant: it, as in model checking, extends the generative range of the testing itself, and (likely) requires a more abstract approach to each mutant than is possible with specific, manually written tests. In short, modifying a test

```
harness cover-harness (H, TARGET)

   H' = H
   for stmt ∈ H':
      if stmt == assert(P):
         stmt = assume(P);
   cover = [
     assume(total_coverage >= TARGET);
     assert(!mutant_covered);]
   insert cover at end of H'.main()
   return H'
```

```
mutant cover-mutant (m)

   n = 0
   m' = m
   for if_stmt c in m':
      if c has no else:
         add [else {}] to c
   for basic_block b in m':
      b = [if !covered[n] {
              covered[n] = 1;
              total_covered += 1;
           }
           b]
      n = n + 1
   for stmt s in m':
      if MUTANT(s):
         s = [{mutant_covered = 1;
               s}]
   m' = [int total_covered = 0;
         int mutant_covered = 0;
         int covered[n];
         m']
   return m'
```

```
trace maxcover (m, H, S, O, U)

   m' = cover-mutant(m)
   T = 0
   trace = ∅
   failed = False
   while (not failed)
      H' = cover-harness(H, T)
      r = scheck(H,m',U(S),O(S))
      if r == VERIFICATION SUCCESSFUL:
        failed = True
      else if r == VERIFICATION FAILED:
        trace = r.trace
        T = trace.read(total_covered) + 1
   return trace
```

Fig. 8 Algorithm 2: Find a maximally covering execution trace that covers a mutant.

```
report check-harness (SUT, M, H, M(H), S, O, U)

   K_H = killed(M, H, S)
   Hkills = ∅; Hequal = ∅; Hbetter = ∅; N = ∅
   for H_i in M(H):
      original = check(H_i, SUT, U(S), O(S))
      if original == VERIFICATION FAILED:
         Hkills += H_i
      else: // check if this kills fewer mutants
         K_{H_i} = killed(M, H, S)
         for k ∈ K_{H_i}:
            if k ∉ K_H: N += (H_i, k)
         if |K_{H_i}| > |K_H|:
            Hbetter += (H_i, K_{H_i})
         if |K_{H_i}| == |K_H|:
            Hequal += (H_i, K_{H_i})
         else:
            Hkills += (H_i, K_{H_i})
   return (Hkills, Hequal, Hbetter, N)
```

Fig. 9 Algorithm 3: Analyze a harness.

harness that generates an essentialy unbounded number of different tests is not, at a high level, very different than modifying a verification harness to detect a mutant. Both operations require not just an analysis of a specific faulty behavior, but an improvement to the *testing process*. The only major difference is that with test generation, the goal is to achieve high probability of detection, since the certainty of verification is lost. In compensation, however, one receives scalability, ease-of-use, and availability in languages without formal verification tools.

Note that we propose mutation testing as a way to *extend* the insight into a test generation effort provided by code coverage [40]. When the generator never produces tests covering a line of code, this can be detected without the expense or difficulty of producing mutants, and interpretation is straightforward: if you are interested in testing the un-covered code, you must modify the generator to cover it. Coverage is good at predicting overall mutation scores [34, 36], but there is a gap between coverage's information and that provided by some surviving mutants. This gap is our interest. Mutants provide insight on the strength of the oracle used to assess generated tests, and the degree to which generated inputs stress it (a line of code can be covered but not thoroughly stressed, since coverage does not consider subtle aspects of, e.g., expression evaluation).

3.1.2 Falsification-Driven Testing

This section explains the adaptations required to apply our approach in the context of automated test generation, rather than formal verification. One aspect of the approach outlined above requires no modifications: examining mutants that are not killed is the same basic process, whether those mutants are not killed by a verification harness or a test generation harness. Similarly, mutating a test harness is not essentially different, though it is less useful in testing than in verification (because of the probabilistic nature of most aggressive testing). Two aspects, however, require some modification.

First, and simplest, the method for generating passing executions is slightly different. Rather than simply negating the specification and adding constraints for code coverage, we modify the testing tool (in our case the TSTL tool for Python [51,47] by adding options to 1) search for a passing test with maximum code coverage and 2) constrain the search to only tests that cover a given statement or branch. These options are: `--trackMaxCoverage <file>`, `--maxMustHitBranch <branch>` and `--maxMustHitStatement <stmt>`.

With this addition to TSTL, finding passing executions to examine is trivial, simply a matter of identifying the location of the mutant. Adding a similar feature to most widely used test generation tools should be relatively easy, given that they work by first generating a test, then executing it and determining its coverage and other properties, such as whether it passes (or generating a test on the fly while determining these properties).

The second change is that the notion of mutant stability changes in two ways. First, the parameter to be optimized is different: while test generation usually requires a maximum test length [7], that parameter is not one with a corresponding computational cost, like a bounded model checking depth. The same test budget can support multiple lengths, and there is no simple correspondence between depth and mutants killed. Mutants killed is monotonic (stable or increasing) in bounded model checking depth; mutants killed is not monotonic in maximum test length [7]. However, there is an analogous parameter in test generation: actual test budget. This parameter, however, only produces *probabilistically* [8] monotonic behavior: one run with budget $X > Y$ may kill fewer mutants than a run with budget Y; however, over a large number of runs, statistically, X-budget tests must outperform, or perform the same as, Y-budget tests.

A key insight is that when estimating how hard a mutant is, a single run that either takes a long time to kill the mutant or a run that fails to kill a (known-killable) mutant is sufficient evidence to assume the mutant is difficult, and helps establish a lower bound on test budget needed to reduce risk of missing faults; in contrast, a single run that quickly kills a mutant does not establish that the mutant is in fact easy: even hard mutants can sometimes be easy to detect. In part we base this idea on empirical evidence (see below), but it can also be justified by a simple theoretical model of test generation.

Unfortunately, the change to a stochastic setting makes determining stability more difficult. It is highly inefficient, once budgets become larger (which is often required in testing), to run each mutant enough times to obtain a good estimate of how long it takes to kill, and doing so requires running already killed mutants, or equivalent mutants, many times. On the other hand, in another sense the problem becomes easier: with bounded model checking, for each mutant m, we have to query wither unwinding depth U is sufficient to kill the mutant, for each U until stability is achieved. Many test generation systems, such as TSTL, support a mode where the tool simply runs until a fault is detected. If we set a large timeout (larger than the largest test budget we are interested in), we can simply run the tool with each mutant and discover how long it takes to kill it (or if it cannot be killed within our maximum possible budget). Finding "stability" then, becomes simply a matter of finding the mutant(s) with the largest time-to-kill, and inspecting those never killed to determine if they are equivalent.

Unfortunately, again due to the stochastic setting, this is not quite sufficient. A hard-to-kill mutant may, by chance, infrequently be detected very quickly; an easy-

Fig. 10 Tests required for detection as difficult of fault/mutant increases.

to-kill mutant may, by chance, sometimes not be detected quickly. Both cases result in inaccurate information, and potentially in a wrong estimate of the correct test budget to use. Fortunately, an asymmetry in the probabilities of these misleading results allows us to proceed without (usually) running each mutant many times.

3.1.3 Estimating Required Budget to Kill a Mutant

The difficulty of detecting a fault (or killing a mutant, which is equivalent) can be simply expressed by a probability of detection, e.g. a trivial fault is detected with almost every generated test (99/100 tests), a typical easy fault is detected frequently (1/100 tests), and a hard fault is detected two orders of magnitude or more less frequently (1/5000 tests). In fact, real faults often seem to fall into such coarse "buckets" of detection ease, though this is not required for our analysis [20].

Figure 10 shows how the minimum number of tests, mean number of tests, and maximum number of tests needed for detection, over 100 runs each consisting of 10,000 tests, vary as the difficulty of a mutant or fault changes. For simplicity, we measure difficulty by increasing N, where the probability of detection is $1/N$. As you can see, while the maximum number of tests required increases rapidly, and the mean number of tests also increases steadily with difficulty, the minimum number of tests required increases very slowly. A concrete example shows why: the chance of getting "lucky" and hitting a test with difficulty 1000 on the first test is indeed small, only 1/1000. However, the chance of getting "unlucky" and

```
float stable-budget (M, k, T, maxTime, equiv)

  budget = 0
  for m ∈ M:
      killTime = -1
      runs = 0
      maxKillTime = -1
      while (killTime < budget) and (runs < k):
          runs = runs + 1
          // Assume killTime is -1 if not killed
          killTime = T(timeout=maxTime)
          if killTime > maxKillTime:
              maxKillTime = killTime
      if maxKillTime == -1:
          equiv = equiv ∪ m
      if maxKillTime > budget:
          budget = killTime
  return budget
```

Fig. 11 Estimating needed budget for mutant-stable testing

failing to detect a test with difficulty 1/10 for even as few as 100 tests is less than 1/30,000.

This means that as soon as a mutant has taken a long time to kill, we can safely assume that it is indeed hard to kill; when a mutant is killed quickly, we run again, based on a tolerance for error in budget estimation (where each run adds large confidence); after k runs we assume a mutant is indeed easy. As soon as one detection takes a long time, or exhausts the budget, we can use that value as an estimate for the difficult of the mutant. While this approach is not perfect, it has the huge advantage of never requiring more than a total runtime slightly more than 1 run with the full test budget, unless k is large. More precisely, to coarsely, but cheaply, estimate a stable test budget for a set of mutants, we use the procedure in Figure 11, which returns a budget and modifies a set of possibly equivalent mutants. The procedure can also be used to set a difficulty for each mutant, in addition to an overall budget, by simply storing the largest `killTime` for that mutant. The procedure takes as input a set of mutants M, a number of trials k, a test procedure T, a maximum test budget to use, $maxTime$, and a (modifiable) structure to store possibly equivalent mutants, $equiv$.

4 Case Studies and Experimental Results

4.1 Algorithm Implementations

Our initial experiments involved relatively small verification problems, based on implementations taken from the web or student code for popular algorithms and data structures. Here we highlight the most interesting of these; we also successfully applied the method to bubble sort and a student's harness for verifying a version of Dijkstra's shortest path algorithm that enables path reconstruction [83]. For the Dijkstra harness, the low mutant kill rate of only 58% showed that while the harness checks incorrect returned paths, it cannot detect when return values

indicate there is no path but one exists. Improving the harness is a substantial exercise, but can be guided by the survival witnesses.

4.1.1 Binary Search

The ideas in this paper grew out of a side project of the first author: to write a follow-up to Jon Bentley's article on verifying binary search [12] in the context of modern software verification tools (and Joshua Bloch's discovery of a bug in the assumptions behind Bentley's proof [14]). The modeling required is moderately complex (to scale well, an abstract "sorted array" that represents all sorted arrays but only introduces variables equal to the number of probes made by the search is essential). In this case, we did not produce an initial, weaker version of the harness, but checked the existing harness using mutants, and determined that all 3 surviving mutants are equivalent to a true binary search.

Checking harness mutants (which took 37 minutes, including computing the kills for the original harness) produced results confirming the belief this is a good harness. Of the 31 compiling harness mutants, 19 failed to verify the correct binary search, and 7 had worse kill rates than the original harness. The remaining 5 harnesses, all with equal kill rates of 86.7%, all modify an assumption to allow the harness to also check size 0 arrays. This doesn't kill any additional mutants, but is harmless as expected. Of the harness mutants with worse kill rates, three are mutants of the assumptions on the nondeterministic value used to make sure that if binary search returns -1, no index in the array actually contains the searched-for item. Two of these mutants are off-by-one-errors that exclude item 0 from the check, an easy-to-make mistake. Both of these fail to kill *exactly* one mutant killed by the correct harness: the mutant that sets the lower bound initially to 1 instead of 0. Traces of passing runs for these mutants show the problem clearly (the sought item at index 0).

4.1.2 Doubly-linked-list Insertion Sort

Another example, making use of recursive data structures and pointer validity checks, is code for inserting an item (in sorted order) into a doubly-linked list [87]. Our initial harness omitted a check for correctness of `prev` pointers. This problem didn't directly prevent mutants from being detected, but pushed the stable size larger, as with the quicksort example above. Looking at a trace of a size 3 run that fails to kill a clearly problematic mutant easily reveals the problem (the results are correct up to `prev` pointers). This example also showed another use of mutants, in that some seemingly problematic surviving mutants actually just showed a pointless redundancy in the implementation, enabling the removal of an entire conditional branch. A harness check (requiring 30 minutes, including computing the mutant kills for the original harness) showed that of the 105 compiling harness mutants, 92 fail to verify the original code. Another 2 have a worse kill rate than the original (which kills 81% of mutants, a low rate due to the code redundancy), and 11 survive. The large number of survivors is due to a redundancy of the final harness, which checks sortedness and the permutation property for both a forward `next` traversal of the list and a `prev` traversal. Omitting any *one* of these (e.g. `prev` sortedness or `next` permutation) the harness can still detect all mutants. Removing

two, however, fails to kill mutants. The two harness mutants with worse kill rates have extremely poor kill rates ($< 50\%$ and $< 25\%$).

4.1.3 AVL Tree

In the case of the AVL tree, the harness we were working with was unable to reach a mutant-stable unwinding without exhausting memory on the verification of the main program (for AVL trees of up to size 5). We are investigating a more efficient harness encoding, based on the inability to reach mutant-stability. Without the notion of mutant-stability, we might have believed the harness was verifying more aspects of the specification than it is able to, at the largest unwinding reached. Unkilled mutants include clearly erroneous behaviors.

4.1.4 Merge With Duplicate Removal

Even a killed mutant—one the harness does detect—can shed critical light on harness vulnerabilities. For example, the code in Figure 12 is a portion of a harness to verify code that merges two sorted arrays, removing all duplicates (the source arrays may contain duplicates or shared items, the output array is guaranteed to be sorted and have all-unique values). This harness detects all non-equivalent mutants of the source code with an unwinding depth of only 2 (the check requires less than a minute).

However, as is well known, many software faults [55] are not represented by a mutant. Because we are model checking, we want our harness to actually rule out *all* bad runs of the program under test. Even a killed mutant's passing executions may show such a problem. Here we see that when the output array's size is 1, the way we have written the duplicate check in fact *assumes* away *all executions*! We check no properties of size 1 output arrays, and a fault that only appears with size = 1 will never be detected. No mutant produces such behavior, but noting an incorrect but passing trace of this run using the CBMC extension lets us see the problem.

4.2 SpiderMonkey Boyer-Moore-Horspool Implementation

Figures 14 and 15 show, respectively, the source code and an initial harness for verification of the Boyer-Moore-Horspool substring finding algorithm [52, 4] in version 1.6 of Mozilla's SpiderMonkey JavaScript engine. Verifying this code presents one immediate issue that is not unusual in verification: how to handle an `assert` in the code being verified. An `assert` at the end of a function or in the main body is typically an additional part of the specification, and is often best left unchanged. An `assert` at the beginning of a function's body, however, is typically a precondition for the code [4]. It is natural to consider changing such an assertion into an `assume` and ignoring any problems produced by calling the code with non-conforming inputs. While this can be a useful technique (for instance when it is hard to write a harness that only produces valid inputs, but easy to filter out the invalid inputs and only verify behavior for those) it is also a dangerous technique. Mutation analysis of the harness shows that 4 is a mutant-stable size (where the same size is used for text length, pattern length, and character set

```
int main () {
  int a[SIZE], b[SIZE], c[SIZE*2];
  int i, v, i1, i2, csize;
  int asize = nondet_int();
  int bsize = nondet_int();
  __CPROVER_assume ((asize >= 0) && (bsize >=0));
  __CPROVER_assume ((asize <= SIZE) && (bsize <= SIZE));
  for (i = 0; i < asize; i++) {
    a[i] = nondet_int();
    __CPROVER_assume((i == 0) || (a[i] >= a[i-1]));
  }
  for (i = 0; i < bsize; i++) {
    b[i] = nondet_int();
    __CPROVER_assume((i == 0) || (b[i] >= b[i-1]));
  }
  csize = merge_sorted_nodups(a, asize, b, bsize, c);
  assert (csize <= (asize + bsize));
  i1 = nondet_int();
  i2 = nondet_int();
  __CPROVER_assume((i1 >= 0) && (i2 >= 0));
  __CPROVER_assume((i1 < csize) & & (i2 < csize));
  __CPROVER_assume(i1 != i2);
  assert(c[i1] != c[i2]);
  v = nondet_int();
  __CPROVER_assume ((v >= 0) && (v < asize));
  v = a[v];
  int found = 0;
  for (i = 0; i < csize; i++) {
    if (c[i] == v)
      found = 1;
  }
  assert (found == 1);
  v = nondet_int();
  __CPROVER_assume ((v >= 0) && (v < bsize));
  v = b[v];
  int found = 0;
  for (i = 0; i < csize; i++) {
    if (c[i] == v)
      found = 1;
  }
  assert (found == 1);
}
```

Fig. 12 Harness for merge_sorted_nodups

size), with a kill rate of 72.3%. On initial examination, the 20 surviving mutants do not seem problematic. A large number involve the JS_ASSERT converted to a __CPROVER_assume, showing the harness cannot tell if the assumption is incorrect, which is not surprising (the harness only generates good inputs, and some of the mutants simply discard too many inputs).

At this point, we were satisfied with our harness, and ran a check on mutants of the harness itself. To our surprise, three mutants of the harness had a better kill rate than the "correct" harness, killing 73.5% of mutants. Investigating these "better" harnesses showed mutants that broke processing of some return values in such a way that, while these harnesses failed to detect certain major bugs in the code, they were able to detect some JS_ASSERT assumption mutants. Guided by

```
int merge_sorted_nodups(int a[], int asize,
                        int b[], int bsize, int c[]) {
  int apos = 0, bpos = 0, cpos = -1, csize = 0;
  while ((apos < asize) || (bpos < bsize)) {
    if ((apos < asize) &&
        ((bpos >= bsize) || (a[apos] < b[bpos]))) {
      if ((cpos == -1) || (c[cpos] != a[apos])) {
        c[++cpos] = a[apos];
        csize++;
      }
      apos++;
    } else {
      if ((cpos == -1) || (c[cpos] != b[bpos])) {
        c[++cpos] = b[bpos];
        csize++;
      }
      bpos++;
    }
  }
  return csize;
}
```

Fig. 13 Code to merge two sorted arrays into one sorted array with no duplicate elements

```
jsint
js_BoyerMooreHorspool(const jschar *text, jsint textlen,
                      const jschar *pat, jsint patlen,
                      jsint start)
{
  jsint i, j, k, m;
  uint8 skip[BMH_CHARSET_SIZE];
  jschar c;
  JS_ASSERT(0 < patlen && patlen <= BMH_PATLEN_MAX);
  for (i = 0; i < BMH_CHARSET_SIZE; i++)
    skip[i] = (uint8)patlen;
  m = patlen - 1;
  for (i = 0; i < m; i++) {
    c = pat[i];
    if (c >= BMH_CHARSET_SIZE)
      return BMH_BAD_PATTERN;
    skip[c] = (uint8)(m - i);
  }
  for (k = start + m;
       k < textlen;
       k += ((c = text[k]) >= BMH_CHARSET_SIZE) ?
            patlen : skip[c]) {
    for (i = k, j = m; ; i--, j--) {
      if (j < 0)
        return i + 1;
      if (text[i] != pat[j])
        break;
    }
  }
  return -1;
}
```

Fig. 14 SpiderMonkey 1.6 Boyer-Moore-Horspool code.

```
#include "bmh.h"
int main() {
  int i;
  unsigned int v;
  char itext[TSIZE];
  char ipat[PSIZE];
  unsigned int itext_s = nondet_uint();
  __CPROVER_assume(itext_s < TSIZE);
  unsigned int ipat_s = nondet_uint();
  __CPROVER_assume(ipat_s < PSIZE);
  printf ("LOG: size text=%u, pat=%u\n", itext_s, ipat_s);
  for (i = 0; i < itext_s; i++) {
    v = nondet_unit();
    __CPROVER_assume((long)v < (long)BMH_CHARSET_SIZE);
    itext[i] = v;
    __CPROVER_assume(itext[i] < BMH_CHARSET_SIZE);
    printf ("LOG: text[%d] = %u\n", i, itext[i]);
  }
  for (i = 0; i < ipat_s; i++) {
    v = nondet_uint();
    __CPROVER_assume((long)v < (long)BMH_CHARSET_SIZE);
    ipat[i] = v;
    __CPROVER_assume(ipat[i] < BMH_CHARSET_SIZE);
    printf ("LOG: pat[%d] = %u\n", i, ipat[i]);
  }
  jsint r = js_BoyerMooreHorspool(itext, itext_s,
                ipat, ipat_s, 0);
  printf ("LOG: return = %d\n", r);
  int pos, ppos, found;
  v = nondet_uint();
  printf ("LOG: looking at %u\n", v);
  __CPROVER_assume(v >= 0);
  if (r == -1) {
    __CPROVER_assume(v < itext_s);
    pos = v; ppos = 0; found = 1;
    while (ppos < ipat_s) {
      printf ("LOG: itext[%d] = %u, ipat[%d] = %u\n",
                pos, itext[pos], ppos, ipat[ppos]);
      if ((pos>=itext_s)||(itext[pos]!=ipat[ppos])) {
        found = 0; break;
      }
      pos++; ppos++;
    }
    assert (!found);
  } else {
    pos = r; ppos = 0;
    while (ppos < ipat_s) {
      assert (itext[pos] == ipat[ppos]);
      pos++; ppos++;
    }
    v = nondet_uint();
    printf ("LOG: looking at %u\n", v);
    __CPROVER_assume(v < r);
    pos = v; ppos = 0; found = 1;
    while (ppos < ipat_s) {
      printf ("LOG: itext[%d] = %u, ipat[%d] = %u\n",
                pos, itext[pos], ppos, ipat[ppos]);
      if ((pos>=itext_s)||(itext[pos]!=ipat[ppos])) {
        found = 0; break;
      }
      pos++; ppos++;
    }
    assert (!found);
  }
}
```

Fig. 15 Boyer-Moore-Horspool harness.

this, we produced a revised harness that raised the kill rate to 79.52%. However, on examining the surviving mutants, we realized that our verification was still unsatisfactory as a good regression for the Boyer-Moore-Horspool code: in particular, if the assertion were ever modified to allow bad inputs to pass through, or otherwise incorrectly changed, we would those bugs. We then changed the `JS_ASSERT` into code that returned a special value to signal assertion failure, and modified the harness once more, allowing some incorrect values to pass through and checking that "assertion failure" happened if, and only if, the inputs were invalid. This harness killed 89.2% of mutants, and the six surviving mutants were easily understood to be equivalent to the BMH code under all valid inputs (in one case we weren't certain about, we had CBMC verify that for all non-assertion violating inputs, this was true up to size 10). The new harness, informed by the harness mutations, in fact had a better mutant kill rate for size 3 (80.7%) than our first harness had at the mutant-stable point. This examples serves as our best evidence of the value of harness mutation.

4.3 Linux Kernel RCU Verification Challenges

Read-Copy-Update (RCU) is a synchronization mechanism sometimes used as a replacement for reader-writer locking for linked structures, allowing extremely lightweight readers [70]. In the limiting case, achieved in server-class builds of the Linux kernel, overhead for entering and exiting an *RCU read-side critical section* (using `rcu_read_lock()` and `rcu_read_unlock()`, respectively) is exactly zero [74], making RCU an excellent choice for read-mostly workloads [70,50,73]. However, lightweight readers mean updaters cannot exclude readers, so updaters must take care to avoid disrupting readers. Updaters typically maintain multiple versions of the portion of the data structure being updated, removing old versions only when readers are no longer accessing them. To this end, RCU provides `synchronize_rcu()`, which waits for a *grace period*: when all pre-existing RCU readers complete. RCU updaters typically remove a data element (rendering it inaccessible to new readers), invoke `synchronize_rcu()`, and then reclaim a removed element.

Because both RCU and the Linux kernel are moving targets, any validation and verification must be both automated and repeatable, for inclusion in a regression-test suite. At present the `rcutorture` stress-test [68] provides some assurance in the form of automated testing, but ideally would be complemented by some formal verification of the implementation(s) in the kernel. An important question is whether available formal verification tools can provide effective additional regression checking for RCU. We use a pair of RCU-related benchmarks [71,72] to provide the beginnings of an answer to this question. The first benchmark applies formal verification to the simplest of the Linux kernel's RCU implementations, Tiny RCU [69], which targets single-CPU systems. This model includes Tiny RCU's handling of idle CPUs as well as its (trivial) grace-period detection scheme. The second benchmark creates the trivial model approximating an RCU implementation for multiprocessor systems shown in Figure 16. In this model, the number of RCU read-side critical sections currently in flight is tracked by the global `rcu_read_nesting_global`, which is atomically incremented by `rcu_read_lock()` and atomically decremented by `rcu_read_unlock()`. This allows `synchronize_rcu()` to

```
 1 static int rcu_read_nesting_global;
 2
 3 static void rcu_read_lock(void)
 4 {
 5   (void)__sync_fetch_and_add(&rcu_read_nesting_global, 2);
 6 }
 7
 8 static void rcu_read_unlock(void)
 9 {
10   (void)__sync_fetch_and_add(&rcu_read_nesting_global, -2);
11 }
12
13 static void synchronize_rcu(void)
14 {
15  for (;;) {
16   if (__sync_fetch_and_xor(&rcu_read_nesting_global, 1)<2)
17     return;
18   SET_NOASSERT();
19   return;
20  }
21 }
```

Fig. 16 Approximate model of RCU

atomically XOR `rcu_read_nesting_global`'s bottom bit to detect whether the current execution has waited for all pre-existing readers (over-approximated by checking the absence of all readers), with `SET_NOASSERT()` being invoked to suppress all future assertions. Although this model has a number of shortcomings, perhaps most prominently excessive read-side ordering, it is capable of detecting common RCU-usage bugs, including failure to wait for an RCU grace period and failure to enclose read-side references in an RCU read-side critical section. Can falsification aid in these two complex, in-progress, verification efforts?

Our efforts are ongoing, due to the complexity of the targeted code (even with support from the primary developer, a co-author of this paper). At this time the investigation of mutants has already provided valuable information about these verifications benchmarks. First, there are two versions of the Tiny RCU verification. The earliest, very preliminary version, kills only 10 of 169 Tiny RCU mutants. Adding code to the harness to account for interrupts in the dyntick-idle handling kills an additional 12 mutants, confirming that the modification increases the strength of the harness. More importantly, the modeling of concurrency in the harness has two versions, one using CBMC support for pthread mutex locks, the other using disabling of assertions to ignore executions that violate locking semantics. The native mutex version allows much faster verification, and catches the original, hand-constructed checks to ensure the harness can detect faults in Tiny RCU. However, the native mutex version fails to kill any mutants, a fact we are currently investigating: without mutants, we would not have been aware of this possibly critical problem, which may be a CBMC bug (in the course of this paper's work, we have uncovered several CBMC bugs) or a harness flaw. In support of the verification, we also generated passing maximal-coverage executions for all mutants of the Tiny RCU code. For 97 of the mutants, there is no passing execution; in many cases, these are not killed: the mutant modifies the concurrency semantics so CBMC has no valid executions to analyze (potentially invalid in some

Fig. 17 Average times for killing/verifying mutants, in seconds.

cases, which must be investigated). For 79 mutants the maximal-coverage passing runs are currently being examined, to determine the best next steps in improving the Tiny RCU harness. For the second benchmark, we have computed mutant kills and find that the kill rates range between 40% and 46%. While these benchmarks are far from complete, and over-simplify the modeling process, they are already able to catch a substantial number of potential RCU usage errors. Again, we have produced passing runs for the surviving mutants to use in enhancing the process. The good news is that while the RCU verification is much more substantial than the above case studies, the time to analyze mutants is not prohibitive. No single model checking run for the Tiny RCU benchmark takes more than 40 seconds, and total runtime for all mutants in the usage benchmarks ranges from just over 12 seconds for a basic litmus test to less than 5 minutes for the most complex of the benchmarks.

4.4 Plausible Verification by Failure to Falsify

A key problem in model checking is the state explosion problem, or, more simply (and more accurately, in that number of states is not always the determining factor in symbolic methods) the problem of scalability. As discussed above, even proving binary search correct over the full domain of integer inputs is not possible within a reasonable time frame. Even when verification is impossible at the desired problem size falsification can provide limited confidence in the correctness of a program. In particular, we observe from all of our experiments that the average time, for any program and harness pair, to verify the original code and all surviving mutants is much higher than the average time to produce a counterexample for killed mutants. Showing that a constraint is satisfiable is, usually, easier than proving it is unsatisfiable. This is not limited to SAT solvers; we used SAT rather than

SMT in our experiments because we generally found Z3 to be slower than CBMC's built in version of MiniSAT[30] in almost all cases, but Z3 also aims to be fast at producing satisfying assignments, not proving UNSAT [75], and our few runs with Z3 showed the same pattern.

Figure 17 shows (with log scales on both axes) the average running times for all experiments (including faulty versions of the harness, incorrect runtime parameters, harness mutation checks, etc.) performed in the course of this work. The general trend is clear: time to verify is usually worse than time to kill, and the worst average time to kill (about 350 seconds) is much better than many average verification times. One use of this relationship is that, in cases where all (non-equivalent) mutants of the SUT are killed, but the SUT verification fails to complete, the SUT might be considered provisionally verified. In particular, the larger the ratio between the timeout for failed verification and the longest kill time for any mutant, the "more likely" to be correct we can consider the SUT (the same holds with respect to memory use limits). This belief can be further justified by modifying the harness to force mutant kills to use large problem sizes, violating the usual inclusiveness rule (that way, if the new size allows a counterexample not previously existing, the mutant killing problem for mutants killable at smaller sizes better approximates the counterexample construction problem for the actual fault).

Additionally, the times shown here (with mean mutant kill time of 16.4 seconds and median mutant kill time of 0.54 seconds) show the general feasibility of the falsification-driven approach. Most of the time, killing mutants is cost-effective. The outliers come from a few difficult problems, arising from buggy harnesses (or harness mutants that resemble buggy harnesses). The much worse cost for surviving mutants is due to a few expensive stubborn mutants: the median verification success time is only 1.5 seconds.

4.5 Automated Test Generation and Falsification

4.5.1 `rcutorture` Case Study

In a first effort to integrate our approach with automated test generation, we applied our method to `rcutorture` [68] itself [2]. While the model checking harnesses we were working with were weak and limited, `rcutorture` is the workhorse of RCU testing, and has resulted in detection of numerous bugs. We generated, again using Andrews' tool, 2,815 valid (compilable) mutants (and 354 invalid mutants). After throwing out all mutants that were equivalent or redundant via TCE [78], we ran `rcutorture` for two minutes on each of the 2,150 remaining mutants. Two minutes is actual testing budget; each mutant took up to 30 minutes to compile in the first place (build failures usually took less time), and a 1-time setup cost of about 30 minutes to produce an image for use in future testing, if we wanted to run `rcutorture` for longer runs (this time without compilation or image-production overhead). Of these 2,150 mutants, only 380 survived the process, yielding an 82% kill rate, showing that `rcutorture` is indeed quite effective at detecting most deviations from the code, and that the code is relatively tightly constrained in behavior.

Manual inspection by the RCU developer (author McKenney) required about 5 minutes per mutant, but with very high variance. Some mutants were immediately dismissed as irrelevant, while other required considerable effort. A good estimate for overall human effort in this case is 25 hours. This is a significant investment of time, but the payoff, in this case, was substantial. In general, the mutants that resulted in changes to kernel code were also mutants that required substantial analysis time.

The results of mutation analysis were five patches to RCU code. Two of these patches were to `rcutorture` code, improving the testing process. Three were patches to RCU code, including one actual bug fix for a bug detected by the changes to `rcutorture`. Another change to RCU was a performance improvement, avoiding the overhead of a `local_irq_save()`/`local_irq_restore` pair, since surviving mutants showed that interrupts were always disabled. A final patch more precisely specified a type, again based on information from mutation testing. Patch details are available at:

- `https://git.kernel.org/pub/scm/linux/kernel/git/tip/tip.git/commit/?id=45fed3e7cfb4001c80cd4bd25249d194a52bfed3`
- `https://git.kernel.org/pub/scm/linux/kernel/git/tip/tip.git/commit/?id=7c9906ca5e582a773fff696975e312cef58a7386`
- `https://git.kernel.org/pub/scm/linux/kernel/git/tip/tip.git/commit/?id=ca1d51ed9809a99d71c23a343b3acd3fd4ad8cbe`
- `https://git.kernel.org/pub/scm/linux/kernel/git/tip/tip.git/commit/?id=6e91f8cb138625be96070b778d9ba71ce520ea7e`
- `https://git.kernel.org/pub/scm/linux/kernel/git/tip/tip.git/commit/?id=f13bad9042dcf9b60b48a0137951b614a2ee24b5`

More details of our `rcutorture` analysis are available in a workshop paper presented at the 2017 International Workshop on Mutation Testing [2].

4.5.2 `pyfakefs` Case Study

We used the a test harness written in the TSTL [47,48,51,49] automated test generation language for testing `pyfakefs` [15] as a second case study in applying our approach to aggressive testing. The `pyfakefs` module is a widely-used tool for Python testing that allows tests to replace real file system usage with the use of a "fake file system" that can contain arbitrary contents, operate faster than a real file system, avoid tests modifying the real file system, and allow fault injection. The harness for `pyfakefs` is a TSTL-based differential testing harness conceptually similar to harnesses used at NASA for file system testing [42,43]. The TSTL tools support a number of testing methods, but in this case study we restricted ourselves to pure random testing.

The muupi [64] mutation generation tool generated 873 mutants of the file `fake_filesystem.py`, the tested part of the `pyfakefs` system. At the time we generated mutants, testing of `pyfakefs` had been going on for a period of months, with many changes and extensions of the test harness, leading to discovery and correction of more than 50 bugs[7] Of the 873 mutants, only 449 (51.4%) were actually

---

[7] See the issues labeled with **TSTL** on the pyfakefs GitHub issue tracker for a history of the testing effort.

Table 1 Easily ignored unkilled mutant categories

Identifier	Explanation	Count
ChangeDiskUsage	Disk usage is hard to model in differential testing	8
usage_change	Disk usage is hard to model in differential testing	2
_last_dev	Differential testing only uses one device	4
_last_ino	inode usage not modeled in differential testing	8
epoch	epoch not modeled in differential testing	3
link_depth	link depth is limited by assumption in testing	4
reversed	reversing lists used as sets or singleton lists is a no-op	5
st_atime	time stat fields not modeled in differential testing	4
st_ctime	time stat fields not modeled in differential testing	4
st_mtime	time stat fields not modeled in differential testing	4
st_mode	mode values differ in ways not modeled in differential testing	9
st_nlink	known discrepancies in nlink behavior to be ignored	7

covered within 30 minutes of testing `pyfakefs`, and so we omitted the mutants of code not covered. Coverage information itself is useful, and showed some (known) limitations of the testing, but, as discussed above, we are interested in the additional information on oracle and input generation strength provided by mutation testing results.

The harness killed 288 of the 449 covered mutants (64.1%), all within 60 seconds. For 279 of these, 5 of 5 attempts to kill within 60 seconds succeeded. For another 7, fewer than 5 (but not less than 3) attempts to kill in 60 seconds succeeded. Extending the test budget to five minutes did not add any newly killed mutants. The mean minimum time to kill a mutant was 0.59 seconds, and the mean maximum time was 3.25 seconds. The longest minimum time to kill a mutant was 44.8 seconds, and the longest mean time to kill a mutant was 42.64 seconds. Our mutation results suggest that 45 seconds of testing is likely to be as effective as 60 seconds (or five minutes) of testing for `pyfakefs`.

Examining the 163 interesting un-killed mutants proved extremely fruitful. First, there were large groups of mutants that could be ignored, as they related to aspects of testing intentionally not performed, such as checking disk usage or atime/ctime/mtime related behavior (these behaviors are ignored due to the problems of differential testing with a reference file system). Additionally, some mutants were obviously uninteresting on inspection (e.g., those that reversed a list where clearly the order of items in a list is not important). Table 1 shows the groups of mutants ignored, and the number of such mutants for each category. Removing these obviously uninteresting mutants left us with 101 mutants to examine. After throwing these mutants out, another two large sets of similar mutants were evident. First, there were 14 mutants that modified a statement containing a check on the fake file system class field indicating whether the file system is case sensitive. We immediately noticed that the harness only produced path components with lowercase characters, meaning that any faults related to handling of case sensitivity would never be detected. Second, there were a large number of mutants referencing methods checking the position of the path separator in a path, and in particular two mutants checking whether the path ends with a path separator. We realized that paths ending in a path separator, or more generally containing extra path separators, were also not being generated, meaning that path normalization related faults would also be missed.

We added these features to the test harness, a simple matter of changing the path component generation line from

```
<component> := <["alpha","beta","gamma","delta", "epsilon","a","b",
    "c","d","e","f","g", "h","omega","lambda","phi"]>
```

to

```
<component> :=<["alpha","beta","gamma","Alpha","Beta","Gamma","a","b",
    "c","d","e","f","g","omega","lambda","phi",""]>
```

and as a result were able to discover the following new faults, almost all of which have since been corrected:

1. https://github.com/jmcgeheeiv/pyfakefs/issues/306
2. https://github.com/jmcgeheeiv/pyfakefs/issues/307
3. https://github.com/jmcgeheeiv/pyfakefs/issues/308
4. https://github.com/jmcgeheeiv/pyfakefs/issues/309
5. https://github.com/jmcgeheeiv/pyfakefs/issues/310
6. https://github.com/jmcgeheeiv/pyfakefs/issues/311
7. https://github.com/jmcgeheeiv/pyfakefs/issues/312
8. https://github.com/jmcgeheeiv/pyfakefs/issues/313
9. https://github.com/jmcgeheeiv/pyfakefs/issues/314
10. https://github.com/jmcgeheeiv/pyfakefs/issues/315
11. https://github.com/jmcgeheeiv/pyfakefs/issues/317
12. https://github.com/jmcgeheeiv/pyfakefs/issues/318
13. https://github.com/jmcgeheeiv/pyfakefs/issues/319
14. https://github.com/jmcgeheeiv/pyfakefs/issues/320
15. https://github.com/jmcgeheeiv/pyfakefs/issues/322

Adding these new features to the test harness and throwing out mutants matching one of our "correctly ignored" classes of mutants, we were able to improve the kill ratio to 81.1% of covered mutants, a result indicating, we believe, a relatively strong oracle, which matches our expectations for a differential testing-based harness (and the very large number of faults thus far detected by the harness).

Examining unkilled mutants using the more advanced, dynamic-analysis-based techniques proposed in this paper was also useful, if more time consuming. For instance, consider a mutant that modifies the line:

```
if (not self.isabs(path)):
```

to

```
if self.isabs(path):
```

This code is easily covered, and stepping through a maximal-coverage test covering it generated using our TSTL extension makes it easy to see why the mutant is not killed. Our harness only generates absolute paths, so the mutant forces the branch to always (instead of never) be taken. That causes execution of the code:

```
path = self.join(getcwd(), path)
```

However, since the current directory is always root ("/") this is a no-op. The change indicates that we can extend our testing by generating relative paths. This means the harness must also be modified to make sure the current directory in both file systems is the same. Unfortunately, without the kind of file system sandboxing that pyfakefs provides, this also changes the current directory in the TSTL test generator itself, breaking the tool. So, while in theory this could be added to testing, in practice the effort required is too large. Discovering this, we can add absolute path queries to our set of ignored mutant types and proceed.

Another mutant introduces a spurious `break` into a loop in a method of the file system's `FakeDirectory` class, `HasParentObject` that checks whether (this is the text of the actual code comment):

```
dir_object is a direct or indirect parent directory, or if both are the same object
```

This code is only called during a `rename` operation, and the change simply makes the function incorrectly return False in cases where discovering the link (that makes a rename invalid) requires traversing multiple levels of indirection. This is clearly possible with our harness, but is likely to be quite difficult, since it requires setting up an invalid rename that is invalid due to at least two levels of indirection. We hypothesized that the test size/search depth limit of 200 operations was making this problem hard to detect, and ran the `tstl` random tester with a depth limit of 500 steps. Within three minutes the mutant was detected. This did not result in a change in the file system harness, but showed that to find all faults in the file system, testing to a more aggressive depth limit is important for thorough testing. And, sure enough, running our harness (without the changes for invalid paths, since not all of the faults thus revealed have been fixed) with depth 1000 revealed what appeared to be a new fault: `https://github.com/jmcgeheeiv/pyfakefs/issues/321`. The test failed when run standalone, without the reference, so it appeared to be a legitimate issue. However, on inspection the test is invalid, since writing to a directory is known to fail. The TSTL harness guards against such a problem, however, so how could the test be generated? It turns out the guard in the harness uses the reference file system (which has a bug, in this case), and incorrectly believes the path involved is not a directory. However, the problem also turns out to be a bug in `pyfakefs` which throws an internal error rather than signaling the correct errno and throwing OSError.

At this point, the utility of examining mutants using our tools to generate witness tests that cover the mutant but pass is quite clear: these were the first two unkilled mutants inspected, chosen at random. The first revealed a desireable but difficult change to the harness (and for now lets us ignore two mutants as clearly unkillable without that change). The second, after investigation, resulted in the discovery of an extremely subtle flaw in the test harness, interacting with a fault in the reference (Mac OS X) file system, and also an actual fault in the tested file system.

5 Discussion: Falsification, Verification, and Popperism

"Those among us who are unwilling to expose their ideas to the hazard of refutation do not take part in the scientific game."
-Popper, *The Logic of Scientific Discovery* [79]

The core idea of this paper is that, while successful verification is the *result* that a developer seeks when verifying a program, it is most meaningful in a context provided by many *failed* verifications. The useful model checking harness (e.g., specification) essentially, is one that prohibits certain execution sequences. This is not controversial; a good property is defined by its rejection of bad behavior. However, in most verification efforts, there is a focus on arriving at a successful verification, which sheds very little light on exactly what has been verified. By focusing on mutants throughout the verification process, our approach shifts the emphasis to one of "verifying" the verification itself by repeatedly *falsifying* claims that various incorrect programs satisfy the property. This is, at a conceptual level, akin to Karl Popper's philosophy of science [79, 80].

For Popper, all scientific knowledge is provisional, and the key to the scientific approach is a critical effort, based on *prohibitive* theories. In brief, Popper proposes that proper science must be strongly grounded in a search for counterexamples. Using mutants as a basis for verification is akin to this approach, with the harness taken to be the "theory" of the empirical behavior of the world. Mutants, in this view, are counterfactual worlds that are likely to violate any correct theory of the actual world. A "scientific theory" (that is, a harness) is proven effective by its ability to be shown to be false in these counterfactual worlds. If we can prove a theory is incorrect for an "incorrect" world and cannot prove it is incorrect for the real world, that gives us greater confidence (always provisional, since our understanding of the world, e.g., any complex software system, is almost always limited and prone to error) that the theory is indeed true of the real world/program. Of course, generating alternative worlds and showing that, for example, special relativity is easily falsified in a world where special relativity does not in fact hold, is not practical in scientific discovery. It is, however, quite easy in the artificial "scientific discovery" sense of verifying properties of computer programs.

Furthermore, many of the theoretical objections to Popper (e.g., such as that we "cannot learn from experience the falsehood of any theory" [60]) do not hold for software correctness problems: we can clearly establish the falsehood of a "theory" in our context by a single counterexample; it is only establishing the truth of a theory, and the value of that truth, that is difficult for us.

The same idea applies to software testing, where there is perhaps even more danger of focusing on a successful result, since small errors in specification are less likely to be detected by lackluster testing efforts. Shifting attention to false claims of correctness, and the ability to detect them, is the mental adjustment, with or without use of program mutants, necessary for good testing and proper attention to not only running the progam but providing an effective oracle [11] for those runs. This point of view both lets us see code coverage [40] as both useful and harmful: code coverage can easily be used as a potentially misleading indicator that the code is "mostly" tested (e.g., "we have 80% coverage, we're done") or as a beneficial guide to code not yet "put through the proving ground" of a test (a very Popperian notion, code that has not been potentially falsified) [1], or a measure of how many times code has been put to the test, with more quantitative coverage counts, such as traditionally provided by `gcov`.

Another way to think about this concept is to note that Popper basically rejects induction, in the sense of drawing general conclusions of truth from particulars (he claims that Hume's famous problem of induction [54] is best solved by stating it cannot be solved). This corresponds to a rejection of one view of testing, where

it is seen as demonstrating that a program works: if we observe enough "good" runs, we can conclude that the program is correct. Dijkstra meant his statement that testing "shows the presence, not the absence of bugs" [19] as a criticism, but Popper implies this is precisely the value of testing, in any context not purely deductive: that is, any context where we are either unable to prove fully that a program satisfies its specification (the usual case) or where we are able to do so, but unsure we really have a complete and perfect specification (which is still almost always the case). In this sense the distinction between testing and verification is not so large, in a Popperian sense.

The novel assumption we make in our use of mutation analysis that goes beyond Popper is our belief that a truer specification likely constrains the programs satisfying it more than a less true specification. Replace "truer" with simply "more scientific" and Popper would likely agree. One additional interesting change is that methods not really (at present) practicable in scientific efforts apply here. For instance, while our notion of tests is *deductive* in that a useful test is one that potentially refutes either the correctness of the program (by failing) or the specification (by allowing a mutant to survive that should not), we can apply random testing. In random testing, most tests are not very useful in a deductive sense: they provide little chance to refute a claim about the program. However, collectively, some randomly generated tests are likely to be powerful for falsification in ways that individual tests designed by humans with the deductive approach in mind seldom achieve. In science the equivalent concept would be to perform a vast set of experiments, with little effort to design them for refutation of a theory, and then scan the data for any results that falsify an existing theory. This seems impractical, to say the least.

The key idea really is that of falsification, in both the enterprise of software correctness and (in Popper's view) the enterprise of scientific, empirical, method. A test, ideally, and a counterexample, by definition, falsifies the claim that a program satisfies some specification. A surviving, non-equivalent, mutant can falsify the claim, less seldom thought about in the act of testing, that a specification or test harness or test suite is, itself, sufficiently falsifiable, and constrains the program enough to be effective. Only repeated, serious effort to falsify all that can be falsified can bring us to (still provisional) confidence. Specifications and test harnesses fulfill the same role in software engineering that proposed natural laws do in the scientific endeavor:

> "Not for nothing do we call the laws of nature 'laws': the more they prohibit the more they say.."
> -Popper, *The Logic of Scientific Discovery* [79]

In this sense, the distinction between specification and verification or testing method is perhaps less important than commonly thought. How these interact to produce a prohibition on bad behavior is shared, and is most critical. Moreover, using Popper's ideas helps us see why we should expect mutation-based falsification to be more fruitful in the context of verification or automated test generation than in manual test construction. A manual test suite is like a scientific theory that consists entirely of basic empirical statements[8] about reality, e.g. that a certain

---

[8] Our terminology here is not quite Popper's, which is somewhat difficult to follow without a lengthy introduction to his classification of statements.

very specific experiment should produce a certain result. While potentially useful in a limited way, and potentially falsifiable, such a "theory" does not even qualify as a theory in Popper's analysis, since he expects theories to make universal statements, and be capable of generating a very large set (or even unbounded set) of such basic empirical statements via deduction: i.e., to produce tests. Falsifying the kind of quasi-theory that a traditional manual test suite represents might be useful, and extend its power, but is a much weaker operation than modifying universal statements with generative power to *produce* a large set of such quasi-theories.

It is not, on the whole, surprising that there should be a correspondence between the ideas of Popper and the effort to verify and test software systems. Popper is clear that his approach is meant as an answer to all the fundamental problems of epistemology. Even such highly practical popular guides to software testing as the well-known book of Kaner, Bach, and Pettichord [56] argue that software testing is essentially an epistemological discipline[9]. We speculate that further close reading of Popper's core works might yield additional insight into software testing, given this epistemological foundation. For example, a notion like that of a *crucial experiment* [79], an experiment devised not just to falsify a given theory, but to guarantee value by falsifying at least one of two competing theories, has a resemblance to a stronger type of differential testing [67], or the methods used in approaches such as regression verification, where the entire goal is to find "experiments" that distinguish two putatively similar systems [85].

5.1 Software Verification and Validation as Science Fiction

Taking the correspondences with Popper's ideas seriously, we arrive at a some high-level conclusions:

- A verification harness (which includes both a specification and a definition of behaviors to apply the specification to) is equivalent to a *scientific theory of the world*.
- A test generation harness (which includes both a specification and a definition of behaviors to apply the specification to) is equivalent to a *scientific theory of the world*.
- Both of these things, then, serve an epistemological purpose: to help understand the actual nature of an external reality (the System Under Test).
- That epistemological purpose is achieved via a route of falsification, or at least serious attempted falsification.
- As in real science, this includes falsification of the theories, via proof that a system does not match the theory.

The science fictional aspect comes in two points. First, mutants allow us, as noted above, to produce hypothetical worlds that are not ours, and make sure that the theory distinguishes our world from these worlds, a new kind of falsification. Second, of course, our purpose is often, unlike science, not purely epistemological in the larger sense. Rather, when we discover our theories do not fit reality, we

---

[9] In fact, Kaner, Bach, and Pettichord explicitly mention Popper, though only in the context of using tests to refute conjectures about the correctness of software, not in the context of attempting to refute the testing effort itself.

often adjust reality to fit our theories. This faces us with a new decision: is our theory inadequate, or is our world inadequate? In both cases, falsification helps us make the decision.

6 Related Work

The idea that a "successful verification" in model checking (or even theorem proving) often simply indicates an inadequate property is long-standing [23,53]. The most recent works in this line of thinking, to our knowledge, use Inductive Validity Cores (IVCs) [31–33] to indicate correspondence between property and constraint on the system under verification.

Use of mutants [13,62] to provide a coverage measure dates back both to these early explorations and relatively recent work [59,9,22]. However, in these efforts the mutation was usually applied to hardware models, and (critically) the surviving mutants were used to, e.g., identify "uncovered" portions of a model, rather than presented to a developer for examination and understanding directly. To our knowledge, no previous work presented passing executions of a source code mutant as a guide to understanding specification weakness. Our modification of the harness is a source-code analogue to attempts to modify logical formulas, e.g., the effort to (in a narrow, vacuity-based sense) produce the strongest passing LTL formula of Chockler et al. [21]. We are not the first to note that model checking, at present, due to the "many obstacles" in proving a system correct, is primarily used for falsification [10]. Most previous work on the topic [10] focused on abstractions based on under-approximation, to ensure counterexamples were not spurious. We instead preserve the goal of verification[10], but drive the verification process, from the human point of view, by repeated falsification of incorrect systems.

More distantly related is the general effort to determine the quality not only of test suites (which is often focused on missing tests within the "range" of testing, not a problem for CBMC) but of test oracles and entire testing infrastructures [11]. The problem of "testing the tester" [38] is fundamental to all efforts to improve software quality. Recent efforts of most interest have focused on measuring *checked coverage* [81,82,76], where a metric tries to make sure the code under test potentially changes the value of an assert, using dynamic slicing [90,86]. This is weaker than requiring the oracle kill a mutant, our goal, but more manageable for testing, where complete behavioral coverage is less feasible than in model checking (and where source code sizes combined with test inadequacy may make hand mutation analysis infeasible).

Our idea of examining successful executions to better understand surviving (and even killed) mutants is a peculiar variation of the fault localization and error explanation problem in model checking [37], with the twist being that we are "explaining" an artificial fault that 1) typically does not cause a test failure (for surviving mutants) and 2) has an obviously known location.

The connection between the ideas of Karl Popper and (software) testing is so obvious that it is fairly commonplace, in both academic and popular work [56]. However, this is almost always in the more narrow connection that a test should

---

[10] Note that we use a model checking approach that already guarantees non-spurious counterexamples, and provides bounded rather than full verification.

try to falsify a program. The link between falsification of specifications or harnesses and "Popperism" appears only in our previous work [39] and in a brief discussion in the direction of the basic idea by Aichernig in his work on model-based mutation testing of reactive systems [3].

According to Mathur [65], the idea of mutation testing itself was first proposed by Lipton, then formalized by DeMillo et al. [28], and practically implemented by Budd et al. [17] in 1980. Mutation analysis subsumes different coverage measures [16,66,77]. Mutants are similar to real faults in terms of both errors produced [27] and difficulty of detection detection [5,6]. Just et al. [55] investigated the relation between mutation score and test case effectiveness using 357 real bugs, and found that the mutation score increased with effectiveness for 75% of cases, which was better than the 46% reported for structural coverage.

7 Conclusions and Future Work

This paper proposes a *falsification-driven* methodology for formal verification and high-quality automated testing, particularly when these tasks are performed by the developers of critical software systems. These developers are usually not experts in formal verification or automated test generation, but in the systems they are verifying. Verification, like testing, we claim, always provisional, in that the potential flaws in our assumptions, specification, and understanding of system behavior tend to leave room for doubt about the correctness of any verification result. Verification of code is not self-explanatory, unlike a counterexample. We propose to take advantage of the use of counterexamples and witnesses and center verification (and testing) around the incorrect programs a verification or test effort fails to prove incorrect. A verification or test effort is considered effective when it finds no faults in the SUT and detects every faulty variation of the SUT. An obvious source of faulty SUT variations is mutants; we also suggest that known-flawed versions of code be included in this set, which all of our tools support, but the key to the method is the generation of a large set of potential buggy versions without additional developer effort.

Given these faulty versions, a developer can examine mutants that a verification effort fails to detect, and (with the algorithms and tools presented in this paper) examine executions showing precisely how a program mutant can "make it through" a verification or testing process without being detected, with assurance that these executions will have high coverage (and thus likely be non-trivial). Developers can also check that a verification or testing harness itself does not have any mutants that 1) verify the SUT while 2) killing more mutants than the original harness. This can help detect very subtle flaws in harnesses, especially those based on bad reasoning about "equivalent" mutants. We demonstrate, as a proof-of-concept, that our approach can be useful for simple but realistic verification efforts, and can contribute to serious systems verification and modeling efforts for complex code such as the Linux kernel RCU implementations. Adapting the approach to testing, it works just as well, actually leading to detection of faults in the RCU implementation and a widely-used Python library.

The bigger picture is that our approach attempts to apply the ideas of Karl Popper's falsification-centered approach to the philosophy of science to the understanding of software systems. In this view, verification is almost always provi-

sional, but we can gain considerable confidence in a verification by making serious attempts to prove its inadequacy.

In future work we plan to continue to apply this falsification-driven approach to the RCU verification, and to other critical systems-software targets, which we expect will lead to discovery of new ways a model checker's ability to ask "what if" questions about program behavior [37, 46] can improve developer understanding of verification efforts. We would also like to integrate falsification-driven verification support into the CBMC Eclipse tools, and use speculative model checking calls and incremental SAT to make mutant analysis available to developers continuously as part of their development/debugging process. Finally, these techniques should also be applicable to verification using, e.g., Java Pathfinder [88].

More fundamentally, it would be useful to perform human studies to determine the actual differences between extending (1) manually produced suites consisting of specific tests with additional specific tests, and (2) extending automated test generation harnesses of the same mutation-killing effectiveness.

References

1. Ahmed, I., Gopinath, R., Brindescu, C., Groce, A., Jensen, C.: Can testedness be effectively measured? In: Proceedings of the 2016 24th ACM SIGSOFT International Symposium on Foundations of Software Engineering, FSE 2016, pp. 547–558. ACM, New York, NY, USA (2016). DOI 10.1145/2950290.2950324. URL http://doi.acm.org/10.1145/2950290.2950324

2. Ahmed, I., Jensen, C., Groce, A., McKenney, P.E.: Applying mutation analysis on kernel test suites: an experience report. In: International Workshop on Mutation Analysis, pp. 110–115 (2017)

3. Aichernig, B.K.: Model-based mutation testing of reactive systems. In: Theories of Programming and Formal Methods, pp. 23–36. Springer (2013)

4. Alipour, M.A., Groce, A., Zhang, C., Sanadaji, A., Caushik, G.: Finding model-checkable needles in large source code haystacks: Modular bug-finding via static analysis and dynamic invariant discovery. In: International Workshop on Constraints in Formal Verification (2013)

5. Andrews, J.H., Briand, L.C., Labiche, Y.: Is mutation an appropriate tool for testing experiments? In: International Conference on Software Engineering, pp. 402–411 (2005)

6. Andrews, J.H., Briand, L.C., Labiche, Y., Namin, A.S.: Using mutation analysis for assessing and comparing testing coverage criteria. IEEE Transactions on Software Engineering **32**(8) (2006)

7. Andrews, J.H., Groce, A., Weston, M., Xu, R.G.: Random test run length and effectiveness. In: Automated Software Engineering, pp. 19–28 (2008)

8. Arcuri, A., Briand, L.: A hitchhiker's guide to statistical tests for assessing randomized algorithms in software engineering. Software Testing, Verification and Reliability **24**(3), 219–250 (2014)

9. Auerbach, G., Copty, F., Paruthi, V.: Formal verification of arbiters using property strengthening and underapproximations. In: Formal Methods in Computer-Aided Design, pp. 21–24 (2010)

10. Ball, T., Kupferman, O., Yorsh, G.: Abstraction for falsification. In: Computer Aided Verification, pp. 67–81 (2005)

11. Barr, E.T., Harman, M., McMinn, P., Shahbaz, M., Yoo, S.: The oracle problem in software testing: A survey. IEEE transactions on software engineering **41**(5), 507–525 (2015)

12. Bentley, J.: Programming pearls: Writing correct programs. Communications of the ACM **26**(12), 1040–1045 (1983)

13. Black, P.E., Okun, V., Yesha, Y.: Mutation of model checker specifications for test generation and evaluation. In: Mutation 2000, pp. 14–20 (2000)

14. Bloch, J.: Extra, extra - read all about it: Nearly all binary searches and mergesorts are broken. `http://googleresearch.blogspot.com/2006/06/extra-extra-read-all-about-it-nearly.html` (2006)
15. mrbean bremen, jmcgeheeiv, et al.: pyfakefs implements a fake file system that mocks the Python file system modules. `https://github.com/jmcgeheeiv/pyfakefs` (2011)
16. Budd, T.A.: Mutation analysis of program test data. Ph.D. thesis, Yale University, New Haven, CT, USA (1980)
17. Budd, T.A., DeMillo, R.A., Lipton, R.J., Sayward, F.G.: Theoretical and empirical studies on using program mutation to test the functional correctness of programs. In: Principles of Programming Languages, pp. 220–233. ACM (1980)
18. Budd, T.A., Lipton, R.J., DeMillo, R.A., Sayward, F.G.: Mutation analysis. Yale University, Department of Computer Science (1979)
19. Buxton, J.N., Randell, B.: Report of a conference sponsored by the NATO science committee. In: NATO Software Engineering Conference, vol. 1969 (1969)
20. Chen, Y., Groce, A., Zhang, C., Wong, W.K., Fern, X., Eide, E., Regehr, J.: Taming compiler fuzzers. In: ACM SIGPLAN Conference on Programming Language Design and Implementation, pp. 197–208 (2013)
21. Chockler, H., Gurfinkel, A., Strichman, O.: Beyond vacuity: Towards the strongest passing formula. In: Proceedings of the 2008 International Conference on Formal Methods in Computer-Aided Design, pp. 24:1–24:8 (2008)
22. Chockler, H., Kroening, D., Purandare, M.: Computing mutation coverage in interpolation-based model checking. IEEE Trans. on CAD of Integrated Circuits and Systems **31**(5), 765–778 (2012)
23. Chockler, H., Kupferman, O., Kurshan, R.P., Vardi, M.Y.: A practical approach to coverage in model checking. In: Computer Aided Verification, pp. 66–78 (2001)
24. Clarke, E., Grumberg, O., McMillan, K., Zhao, X.: Efficient generation of counterexamples and witnesses in symbolic model checking. In: Design Automation Conference, pp. 427–432 (1995)
25. Clarke, E.M., Grumberg, O., Peled, D.: Model Checking. MIT Press (2000)
26. Cuoq, P., Monate, B., Pacalet, A., Prevosto, V., Regehr, J., Yakobowski, B., Yang, X.: Testing static analyzers with randomly generated programs. In: NASA Formal Methods Symposium, pp. 120–125 (2012)
27. Daran, M., Thévenod-Fosse, P.: Software error analysis: A real case study involving real faults and mutations. In: ACM SIGSOFT International Symposium on Software Testing and Analysis, pp. 158–171. ACM (1996)
28. DeMillo, R.A., Lipton, R.J., Sayward, F.G.: Hints on test data selection: Help for the practicing programmer. Computer **4**(11) (1978)
29. Dijkstra, E.W.: A Discipline of Programming. Prentice-Hall, Englewood Cliffs, New Jersey (1976)
30. Een, N., Sorensson, N.: An extensible SAT-solver. In: Symposium on the Theory and Applications of Satisfiability Testing (SAT), pp. 502–518 (2003)
31. Ghassabani, E., Gacek, A., Whalen, M.W.: Efficient generation of inductive validity cores for safety properties. In: ACM SIGSOFT Symposium on the Foundations of Software Engineering, pp. 314–325 (2016)
32. Ghassabani, E., Gacek, A., Whalen, M.W., Heimdahl, M.P.E., Wagner, L.G.: Proof-based coverage metrics for formal verification. In: IEEE/ACM International Conference on Automated Software Engineering, pp. 194–199 (2017)
33. Ghassabani, E., Whalen, M.W., Gacek, A.: Efficient generation of all minimal inductive validity cores. In: FMCAD, pp. 31–38 (2017)
34. Gligoric, M., Groce, A., Zhang, C., Sharma, R., Alipour, A., Marinov, D.: Comparing non-adequate test suites using coverage criteria. In: International Symposium on Software Testing and Analysis, pp. 302–313 (2013)
35. Gligoric, M., Groce, A., Zhang, C., Sharma, R., Alipour, A., Marinov, D.: Guidelines for coverage-based comparisons of non-adequate test suites. ACM Transactions on Software Engineering and Methodology (accepted for publication)
36. Gopinath, R., Jensen, C., Groce, A.: Code coverage for suite evaluation by developers. In: International Conference on Software Engineering, pp. 72–82 (2014)
37. Groce, A.: Error explanation with distance metrics. In: Tools and Algorithms for the Construction and Analysis of Systems, pp. 108–122 (2004)
38. Groce, A.: (Quickly) testing the tester via path coverage. In: Workshop on Dynamic Analysis (2009)

39. Groce, A., Ahmed, I., Jensen, C., McKenney, P.E.: How verified is my code? falsification-driven verification. In: 30th IEEE/ACM International Conference on Automated Software Engineering, ASE 2015, Lincoln, NE, USA, November 9-13, 2015, pp. 737–748 (2015). DOI 10.1109/ASE.2015.40
40. Groce, A., Alipour, M.A., Gopinath, R.: Coverage and its discontents. In: Onward! Essays, pp. 255–268 (2014)
41. Groce, A., Erwig, M.: Finding common ground: choose, assert, and assume. In: Workshop on Dynamic Analysis, pp. 12–17 (2012)
42. Groce, A., Holzmann, G., Joshi, R.: Randomized differential testing as a prelude to formal verification. In: International Conference on Software Engineering, pp. 621–631 (2007)
43. Groce, A., Holzmann, G., Joshi, R., Xu, R.G.: Putting flight software through the paces with testing, model checking, and constraint-solving. In: Workshop on Constraints in Formal Verification, pp. 1–15 (2008)
44. Groce, A., Joshi, R.: Exploiting traces in program analysis. In: Tools and Algorithms for the Construction and Analysis of Systems, pp. 379–393 (2006)
45. Groce, A., Joshi, R.: Random testing and model checking: Building a common framework for nondeterministic exploration. In: Workshop on Dynamic Analysis, pp. 22–28 (2008)
46. Groce, A., Kroening, D.: Making the most of BMC counterexamples. Electron. Notes Theor. Comput. Sci. **119**(2), 67–81 (2005)
47. Groce, A., Pinto, J.: A little language for testing. In: NASA Formal Methods Symposium, pp. 204–218 (2015)
48. Groce, A., Pinto, J., Azimi, P., Mittal, P.: TSTL: a language and tool for testing (demo). In: ACM International Symposium on Software Testing and Analysis, pp. 414–417 (2015)
49. Groce, A., Pinto, J., Azimi, P., Mittal, P., Holmes, J., Kellar, K.: TSTL: the template scripting testing language. `https://github.com/agroce/tstl` (2015)
50. Guniguntala, D., McKenney, P.E., Triplett, J., Walpole, J.: The read-copy-update mechanism for supporting real-time applications on shared-memory multiprocessor systems with Linux. IBM Systems Journal **47**(2), 221–236 (2008)
51. Holmes, J., Groce, A., Pinto, J., Mittal, P., Azimi, P., Kellar, K., O'Brien, J.: TSTL: the template scripting testing language. International Journal on Software Tools for Technology Transfer (2016). DOI \url{https://doi.org/10.1007/s10009-016-0445-y}. Online first
52. Horspool, R.N.: Practical fast searching in strings. Software - Practice & Experience **10**(6), 501–506 (1980)
53. Hoskote, Y., Kam, T., Ho, P.H., Zhao, X.: Coverage estimation for symbolic model checking. In: ACM/IEEE Design Automation Conference, pp. 300–305 (1999)
54. Hume, D.: An Enquiry Concerning Human Understanding. London (1748)
55. Just, R., Jalali, D., Inozemtseva, L., Ernst, M.D., Holmes, R., Fraser, G.: Are mutants a valid substitute for real faults in software testing? In: ACM SIGSOFT Symposium on Foundations of Software Engineering, pp. 654–665 (2014)
56. Kaner, C., Bach, J., Pettichord, B.: Lessons Learned in Software Testing: a Context-Driven Approach. Wiley (2001)
57. Kroening, D., Clarke, E.M., Lerda, F.: A tool for checking ANSI-C programs. In: Tools and Algorithms for the Construction and Analysis of Systems, pp. 168–176 (2004)
58. Kroening, D., Strichman, O.: Efficient computation of recurrence diameters. In: Verification, Model Checking, and Abstract Interpretation, pp. 298–309 (2003)
59. Kupferman, O., Li, W., Seshia, S.: A theory of mutations with applications to vacuity, coverage, and fault tolerance. In: Formal Methods in Computer-Aided Design, pp. 1–9 (2008)
60. Lakatos, I.: The role of crucial experiments in science. Studies in History and Philosophy of Science Part A **4**(4), 309–325 (1974)
61. Lawlor, R.: quicksort.c. `http://www.comp.dit.ie/rlawlor/Alg_DS/sorting/quickSort.c`. Referenced April 20, 2015
62. Lee, T.C., Hsiung, P.A.: Mutation coverage estimation for model checking. In: Automated Technology for Verification and Analysis, pp. 354–368 (2004)
63. Lipton, R.J.: Fault diagnosis of computer programs. Tech. rep., Carnegie Mellon Univ. (1971)
64. Liu, X.: muupi mutation tool. `https://github.com/apepkuss/muupi` (2016)
65. Mathur, A.P.: Foundations of Software Testing. Addison-Wesley (2012)
66. Mathur, A.P., Wong, W.E.: An empirical comparison of data flow and mutation-based test adequacy criteria. Journal of Software Testing, Verification and Reliability **4**(1), 9–31 (1994)

67. McKeeman, W.: Differential testing for software. Digital Technical Journal of Digital Equipment Corporation **10(1)**, 100–107 (1998)
68. McKenney, P.E.: RCU torture test operation. `https://www.kernel.org/doc/Documentation/RCU/torture.txt`
69. McKenney, P.E.: Re: [PATCH fyi] RCU: the bloatwatch edition (2009). Available: `http://lkml.org/lkml/2009/1/14/449` [Viewed January 15, 2009]
70. McKenney, P.E.: Structured deferral: synchronization via procrastination. Commun. ACM **56**(7), 40–49 (2013). DOI 10.1145/2483852.2483867. URL `http://doi.acm.org/10.1145/2483852.2483867`
71. McKenney, P.E.: Verification challenge 4: Tiny RCU. `http://paulmck.livejournal.com/39343.html` (2015)
72. McKenney, P.E.: Verification challenge 5: Uses of RCU. `http://paulmck.livejournal.com/39793.html` (2015)
73. McKenney, P.E., Eggemann, D., Randhawa, R.: Improving energy efficiency on asymmetric multiprocessing systems (2013). `https://www.usenix.org/system/files/hotpar13-poster8-mckenney.pdf`
74. McKenney, P.E., Slingwine, J.D.: Read-copy update: Using execution history to solve concurrency problems. In: Parallel and Distributed Computing and Systems, pp. 509–518. Las Vegas, NV (1998)
75. de Moura, L.M., Bjørner, N.: Z3: an efficient SMT solver. In: Tools and Algorithms for the Construction and Analysis of Systems, pp. 337–340 (2008)
76. Murugesan, A., Whalen, M.W., Rungta, N., Tkachuk, O., Person, S., Heimdahl, M.P.E., You, D.: Are we there yet? determining the adequacy of formalized requirements and test suites. In: NASA Formal Methods Symposium, pp. 279–294 (2015)
77. Offutt, A.J., Voas, J.M.: Subsumption of condition coverage techniques by mutation testing. Tech. rep., Technical Report ISSE-TR-96-01, Information and Software Systems Engineering, George Mason University (1996)
78. Papadakis, M., Jia, Y., Harman, M., Traon, Y.L.: Trivial compiler equivalence: A large scale empirical study of a simple fast and effective equivalent mutant detection technique. In: International Conference on Software Engineering (2015)
79. Popper, K.: The Logic of Scientific Discovery. Hutchinson (1959)
80. Popper, K.: Conjectures and Refutations: The Growth of Scientific Knowledge (1963)
81. Schuler, D., Zeller, A.: Assessing oracle quality with checked coverage. In: International Conference on Software Testing, Verification and Validation, pp. 90–99 (2011)
82. Schuler, D., Zeller, A.: Checked coverage: an indicator for oracle quality. Software Testing, Verification, and Reliability **23**(7), 531–551 (2013)
83. scvalex: Finding all paths of minimum length to a node using dijkstras algorithm. `https://compprog.wordpress.com/2008/01/17/finding-all-paths-of-minimum-length-to-a-node-using-dijkstras-algorithm/` (2008)
84. Stout, R.: If Death Ever Slept. Viking (1957)
85. Strichman, O., Godlin, B.: Regression verification-a practical way to verify programs. Verified Software: Theories, Tools, Experiments pp. 496–501 (2008)
86. Tip, F.: A survey of program slicing techniques. Journal of programming languages **3**, 121–189 (1995)
87. visar: [SOLVED] doubly linked list insertion sort in C. `http://www.linuxquestions.org/questions/programming-9/doubly-linked-list-insertion-sort-in-c-4175415860/` (2012)
88. Visser, W., Havelund, K., Brat, G., Park, S., Lerda, F.: Model checking programs. Automated Software Engineering **10**(2), 203–232 (2003)
89. Yang, X., Chen, Y., Eide, E., Regehr, J.: Finding and understanding bugs in C compilers. In: ACM SIGPLAN Conference on Programming Language Design and Implementation, pp. 283–294 (2011)
90. Zhang, X., Gupta, R., Zhang, Y.: Precise dynamic slicing algorithms. In: International Conference on Software Engineering, pp. 319–329 (2003)